

# **EECS 489**

# **Computer Networks**

**Winter 2024**

Mosharaf Chowdhury

*Material with thanks to Aditya Akella, Sugih Jamin, Philip Levis, Sylvia Ratnasamy, Peter Steenkiste, and many other colleagues.*

# Recap: End-to-end principle

---

- Functions that can be *completely* and *correctly* implemented *only* with the knowledge of application end host, should not be pushed into the network
- **Fate sharing**: fail together or don't fail at all
  - “it is acceptable to lose the state information associated with an entity if, at the same time, the entity itself is lost”

# Agenda

---

- HTTP and the Web
- Improving HTTP Performance

# The Web: Precursor

---

- 1945, [Vannevar Bush](#), Memex
  - Concept of the web based on microfilms
- 1967, [Ted Nelson](#), Project Xanadu
  - A world-wide publishing network to store information as connected literature
  - Coined the term “Hypertext”
- 1968, [Douglas Engelbart](#), NLS (oN-Line System)
  - The mother of all demos

# The Web: History

---

- World Wide Web (WWW): a distributed database of “pages” linked through Hypertext Transport Protocol (HTTP)
  - First HTTP implementation – 1990
    - » **Tim Berners-Lee** at CERN
  - HTTP/0.9 – 1991
    - » Simple GET command for the Web
  - HTTP/1.0 – 1992
    - » Client/server information, simple caching



AWARD WINNER

Sir Tim Berners-Lee

ACM A. M. Turing Award (2016)

ACM Software System Award (1995)

2016 ACM A.M. Turing Award

# The Web: History (cont'd)

---

- World Wide Web (WWW): a distributed database of “pages” linked through Hypertext Transport Protocol (HTTP)
  - HTTP/1.1 – 1996
    - » Performance and security optimizations
  - HTTP/2 – 2015
    - » Latency optimizations via request multiplexing over single TCP connection
    - » Binary protocol instead of text
    - » Server push

# The Web: History (cont'd)

---

- World Wide Web (WWW): a distributed database of “pages” linked through Hypertext Transport Protocol (HTTP)
  - HTTP/3 – 2022
    - » Built on top of QUIC, which is a user-space congestion control protocol on top of UDP
    - » Solves the **head-of-line (HOL) blocking problem** when multiplexing over a single TCP connection

# What does it consist of?

---

- Who uses it?
- Who provides the content?
- How do they communicate?
  
- How do we find the content?
- How is the content organized?
- How is it displayed?



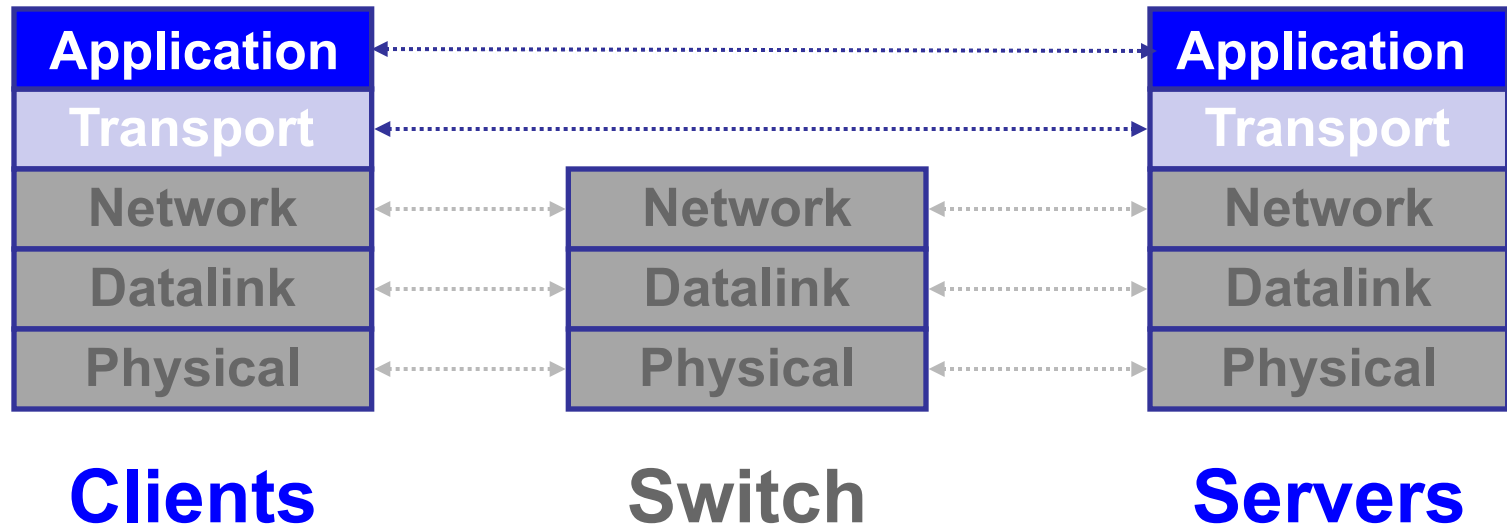
# Web components

---

- Infrastructure:
  - Clients
  - Servers (DNS, CDN, Datacenters)
- Content:
  - URL: naming content
  - HTML: formatting content
- Protocol for exchanging information: **HTTP**

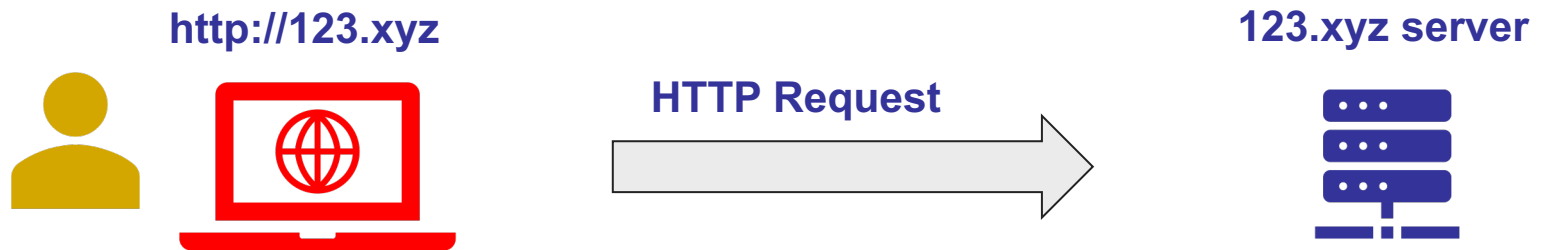
# Why is there nothing about the network?

---



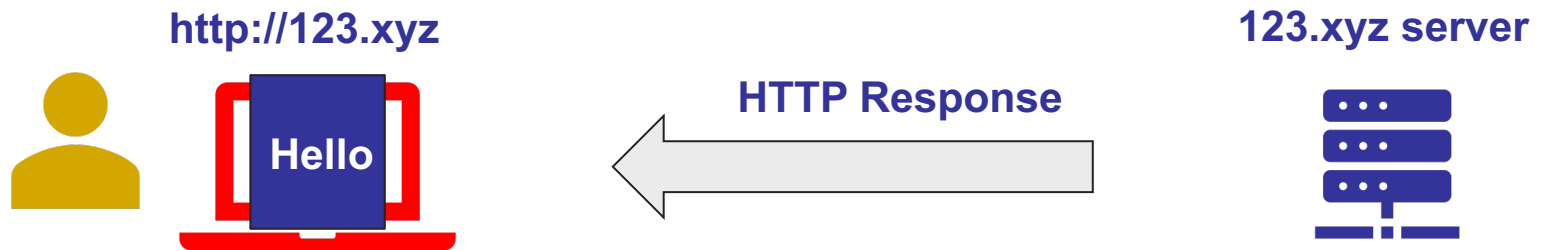
# What we want

---



# What we get

---



# URL: Uniform Record Locator

---

- `protocol://host-name[:port]/directory-path/resource`
- Extend the idea of hierarchical hostnames to include anything in a file system
  - `https://github.com/mosharaf/eecs489/blob/f24/slides/012224.pptx`
- Extend to program executions as well...
  - `https://www.google.com/search?q=eecs489`
  - Server-side processing can be included in the name

# URL: Uniform Record Locator

---

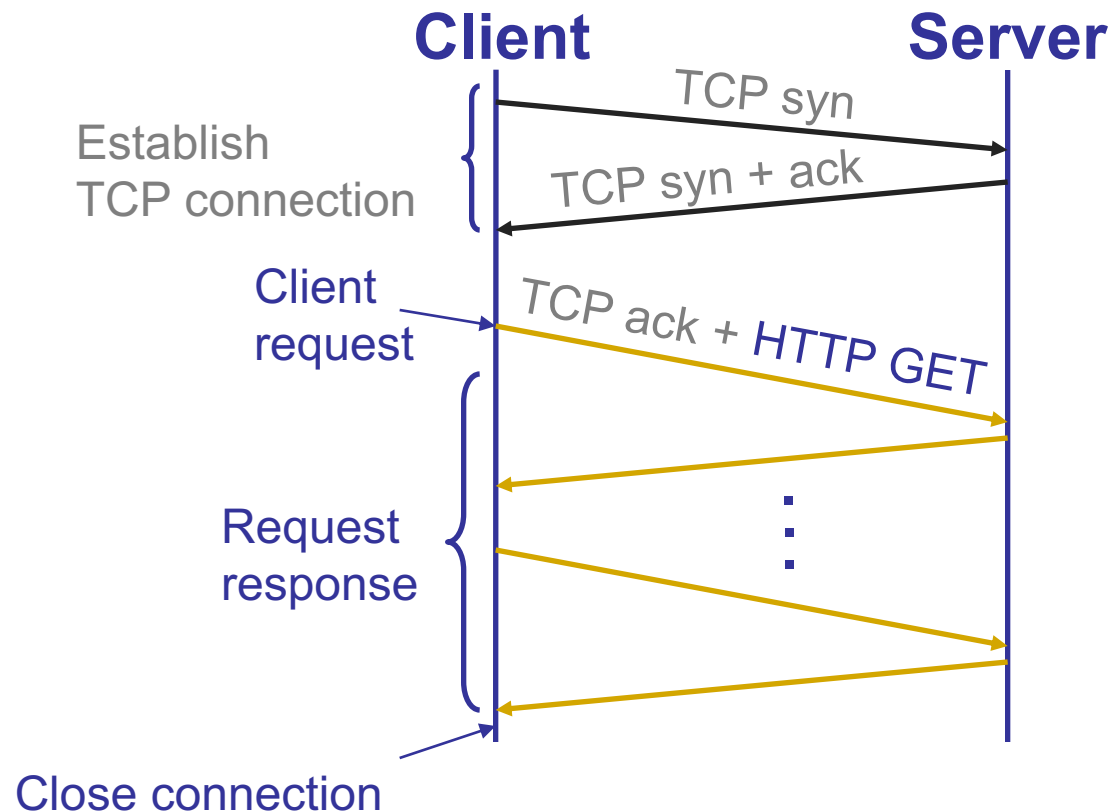
- `protocol://host-name[:port]/directory-path/resource`
  - `protocol`: http, ftp, https, smtp, rtsp, *etc.*
  - `host-name`: DNS name, IP address
  - `port`: defaults to protocol's standard port
    - » *E.g.*, http: 80, https: 443
  - `directory path`: hierarchical, reflecting file system
  - `resource`: Identifies the desired resource

# Hyper Text Transfer Protocol (HTTP)

---

- Client-server architecture
  - Server is “always on” and “well known”
  - Clients initiate contact to server
- Synchronous request/reply protocol
  - Runs over TCP, Port 80
- Stateless
- ASCII format before HTTP/2 but Binary since then

# Steps in canonical HTTP request/response





# Method types (HTTP 1.1)

---

- GET, HEAD
- POST
  - Send information (e.g., web forms)
- PUT
  - Uploads file in entity body to path specified in URL field
- DELETE
  - Deletes file specified in the URL field

# Client-to-server communication

- HTTP Request Message

- Request line: method, resource, and protocol version

*request line* → GET /somedir/page.html HTTP/1.1

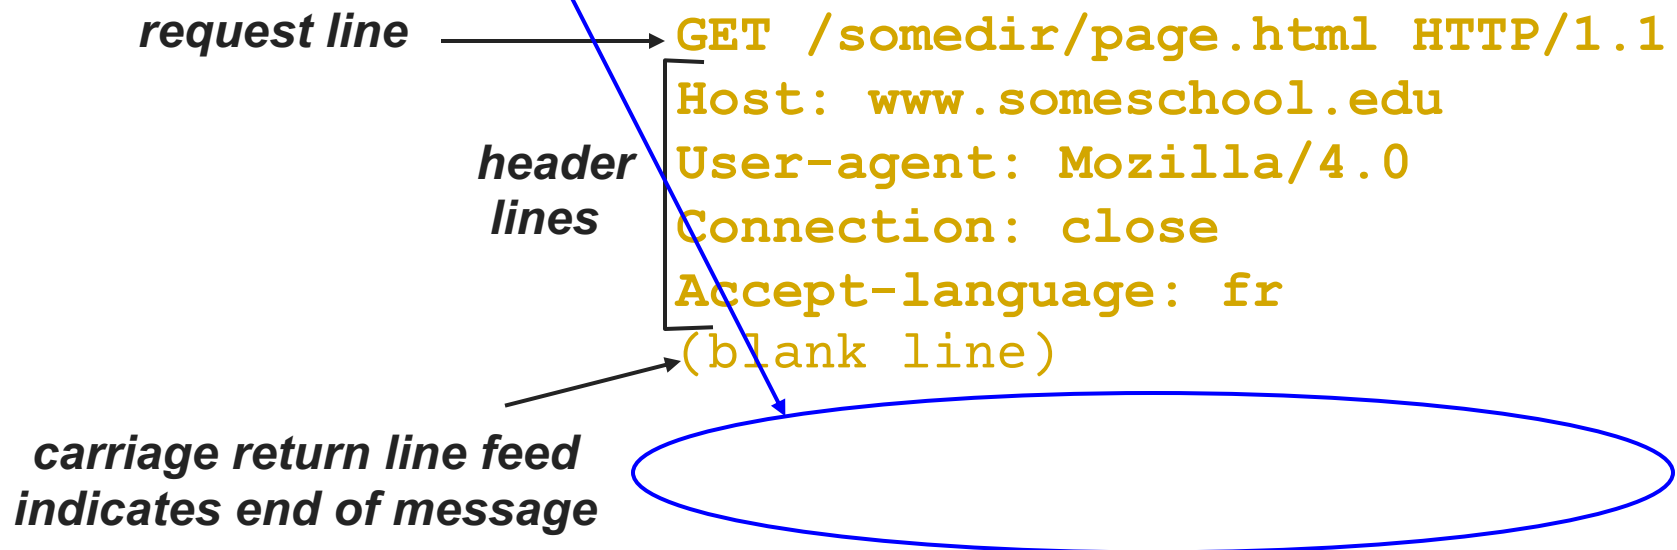
*header lines* → Host: www.someschool.edu  
User-agent: Mozilla/4.0  
Connection: close  
Accept-language: fr  
(blank line)

*carriage return line feed  
indicates end of message*

# Client-to-server communication

- HTTP Request Message

- Request line: method, resource, and protocol version
- Request headers: provide info or modify request
- Body: optional data (e.g., to “POST” data to server)



# Server-to-client communication

- HTTP Response Message

- Status line: protocol version, status code, status phrase
- Response headers: provide information
- Body: optional data

**status line**

(protocol, status code, status phrase)

**header lines**

**data**

HTTP/1.1 200 OK

Connection close

Date: Thu, 06 Jan 2017 12:00:15 GMT

Server: Apache/1.3.0 (Unix)

Last-Modified: Mon, 22 Jun 2006 ...

Content-Length: 6821

Content-Type: text/html

(blank line)

data data data data data ...

e.g., requested HTML file

# HTTP is stateless

---

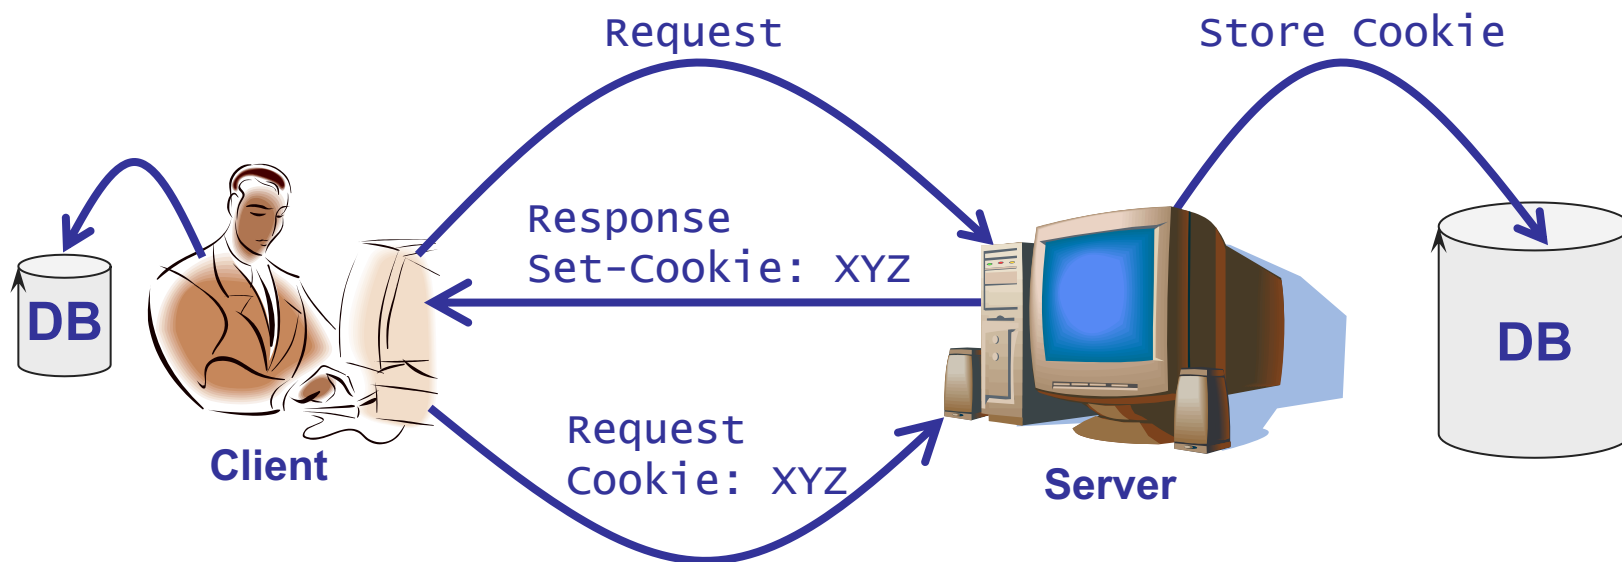
- Each request-response treated **independently**
  - Servers not required to retain state
- **Good**: Improves scalability on the server-side
  - Failure handling is easier
  - Can handle higher rate of requests
  - Order of requests doesn't matter
- **Bad**: Some applications need persistent state
  - Need to uniquely identify user or store temporary info
  - e.g., Shopping cart, user profiles, usage tracking, ...

# How does a stateless protocol keep state?

---

# State in a stateless protocol: Cookies

- **Client-side** state maintenance
  - Client stores small state on behalf of server
  - Client sends state in future requests to the server
- Can provide authentication



# Beyond cookies

---

- Cookies provide excellent marketing opportunities and create concerns for privacy
  - Advertising companies tracks your preferences and viewing history across sites
- Many are trying to curtail the (mis)use of cookies
  - Example: Google Chrome is deprecating third-party cookies in 2024



---

**5-MINUTE BREAK!**

# Announcements

---

- Assignment 1 is due on Jan 26, 2024
- Group formation for A2-A4 by Jan 27 2024
  - [https://docs.google.com/forms/d/e/1FAIpQLSfnVI9BINtZL07Btlb4BGhQm5V4rvQbgXAsGgsYH2TfJroHN/g/viewform?usp=sf\\_link](https://docs.google.com/forms/d/e/1FAIpQLSfnVI9BINtZL07Btlb4BGhQm5V4rvQbgXAsGgsYH2TfJroHN/g/viewform?usp=sf_link)

# Performance goals

---

- User
  - Fast downloads (not identical to low-latency communication!)
  - High availability
- Content provider
  - Happy users (hence, above)
  - Cost-effective infrastructure
- Network (secondary)
  - Avoid overload

# Solutions?

Improve networking protocols including HTTP, TCP, etc.

- User
  - Fast downloads (not identical to low-latency communication!)
  - High availability
- Content provider
  - Happy users (hence, above)
  - Cost-effective infrastructure
- Network (secondary)
  - Avoid overload

# Solutions?

Improve networking protocols including HTTP, TCP, etc.

- User
  - Fast downloads (not identical to low-latency communication!)
  - High availability
- Content provider
  - Happy users (hence, above)
  - Cost-effective infrastructure
- Network (secondary)
  - Avoid overload

Caching and replication

# Solutions?

Improve networking protocols including HTTP, TCP, etc.

- User
  - Fast downloads (not identical to low-latency communication!)
  - High availability
- Content provider
  - Happy users (hence, above)
  - Cost-effective infrastructure
- Network (secondary)
  - Avoid overload

Caching and replication

Exploit economies of scale; e.g., webhosting, CDNs, datacenters

# HTTP performance

---

- Most Web pages have multiple objects
  - e.g., HTML file and a bunch of embedded images
- How do you retrieve those objects (naively)?
  - One item at a time
- New TCP connection per (small) object!


# Object request response time

- RTT (round-trip time)

- Time for a small packet to travel from client to server and back

- Response time

- 1 RTT for TCP setup
- 1 RTT for HTTP request and first few bytes
- Transmission time
- **Total** = 2RTT + Transmission Time





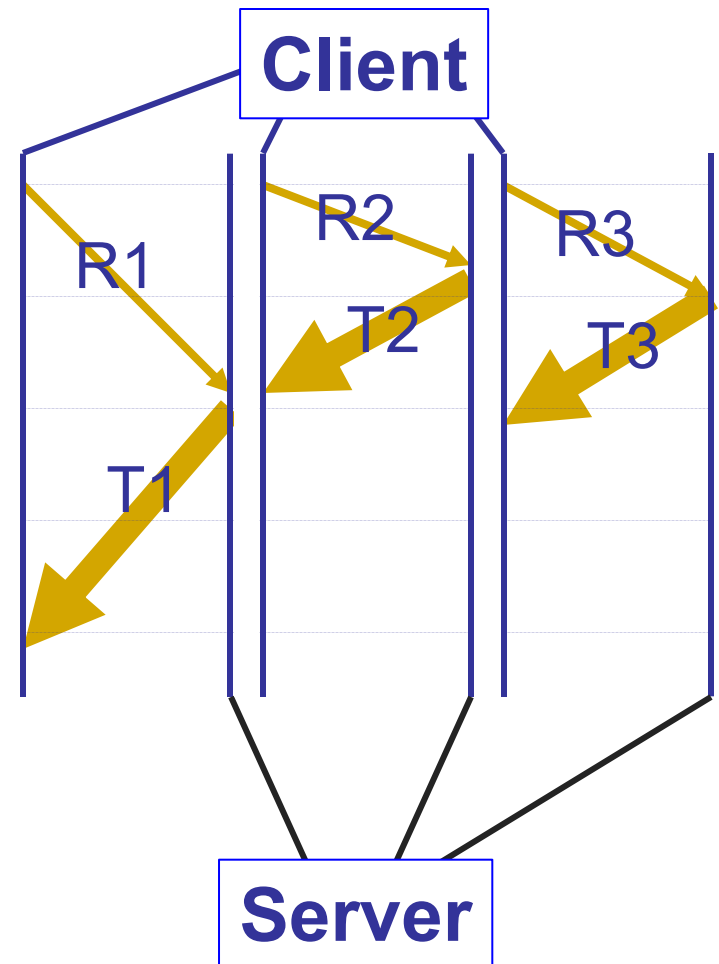
# Non-persistent connections

---

- Default in HTTP/1.0
- $2RTT + \Delta$  for each object in the HTML file!
  - One more  $2RTT + \Delta$  for the HTML file itself
- Doing the same thing over and over again
  - Inefficient

# Concurrent requests and responses

- Use multiple connections in parallel
  - Does not necessarily maintain order of responses
- 
- Client = 😊
  - Content provider = 😊
  - Network = 😞 Why?



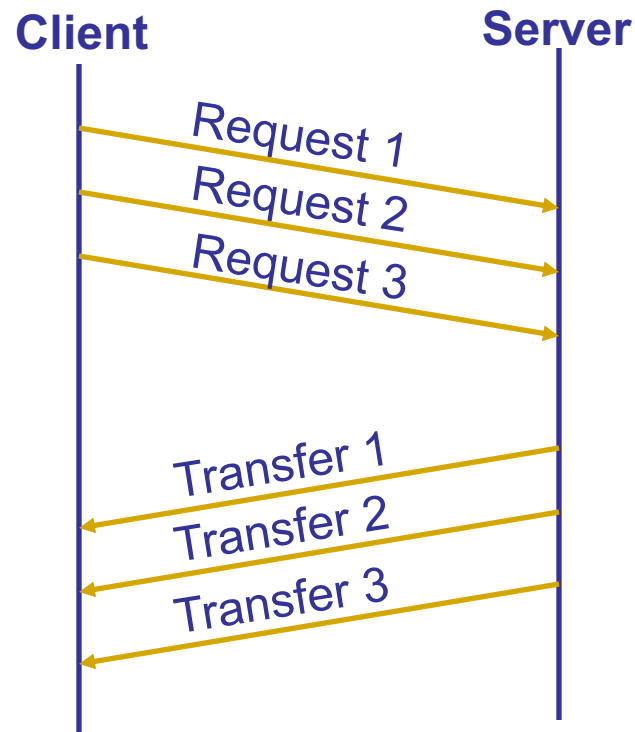
# Persistent connections

---

- Maintain TCP connection across multiple requests
  - Including transfers subsequent to current page
  - Client or server can tear down connection
- Advantages
  - Avoid overhead of connection set-up and tear-down
  - Allow underlying layers (e.g., TCP) to learn about RTT and bandwidth characteristics
- Default in HTTP/1.1

# Pipelined requests & responses

- Batch requests and responses to reduce the number of packets
  - Multiple requests can be contained in one TCP segment
- Data are sent in a FIFO manner
  - Can lead to **head-of-line (HOL) blocking** if many small responses follow a large one
  - Not supported by default by major browsers circa 2015
- **Solution**
  - Priority and preemption



# Scorecard: Getting $n$ small objects

---

- Time dominated by latency
- One-at-a-time:  $\sim 2n$  RTT
- $m$  concurrent:  $\sim 2\lceil n/m \rceil$  RTT
- Persistent:  $\sim (n+1)$  RTT
- Pipelined:  $\sim 2$  RTT
- Pipelined and Persistent:  $\sim 2$  RTT first time; RTT later for another  $n$  from the same site

# Scorecard: Getting $n$ large objects each of size $F$

---

- Time dominated by TCP throughput  $B_C$  ( $\leq B_L$ ), where bottleneck link bandwidth is  $B_L$ 
  - Assuming all TCP connections go through the same bottleneck link
- One-at-a-time:  $\sim nF/B_C$
- $m$  concurrent:  $\sim nF/(mB'_C)$ 
  - Assuming each TCP connection gets the same throughput ( $B'_C$ ), where  $mB'_C \leq B_L$
- Pipelined and/or persistent:  $\sim nF/B_C$ 
  - The only thing that helps is higher throughput

# Caching

---

- Why does caching work?
  - Exploits locality of reference
- How well does caching work?
  - Very well, up to a limit
  - Large overlap in content
  - But many unique requests
    - » A universal story!
    - » Effectiveness of caching grows logarithmically with size

# Caching: How

---

- Modifier to GET requests:
  - **If-modified-since** – returns “not modified” if resource not modified since specified time

```
GET /somedir/page.html HTTP/1.1
Host: www.someschool.edu
User-agent: Mozilla/4.0
If-modified-since: wed, 18 Jan 2017 10:25:50 GMT
(blank line)
```



# Caching: How

---

- Modifier to GET requests:
  - `If-modified-since` – returns “not modified” if resource not modified since specified time
- Client specifies “`if-modified-since`” time in request
- Server compares this against “last modified” time of resource
- Server returns “Not Modified” if resource has not changed
- .... or a “OK” with the latest version otherwise

# Caching: How

---

- Modifier to GET requests:
  - **If-modified-since** – returns “not modified” if resource not modified since specified time
- Response header:
  - **Expires** – how long it's safe to cache the resource
  - **No-cache** – ignore all caches; always get resource directly from server

# Caching: Where?

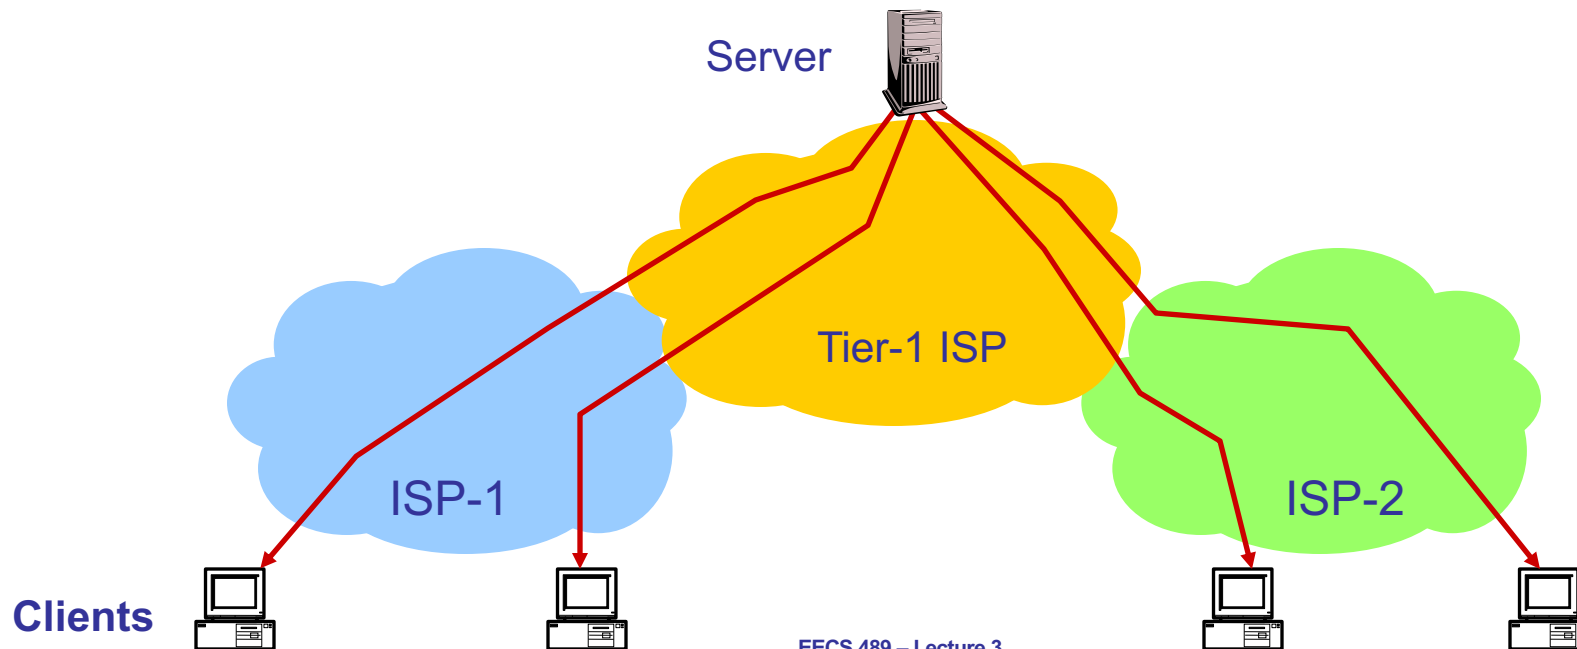
---

- Options
  - Client (browser)
  - Forward proxies
  - Reverse proxies
  - Content Distribution Network

# Caching: Where?

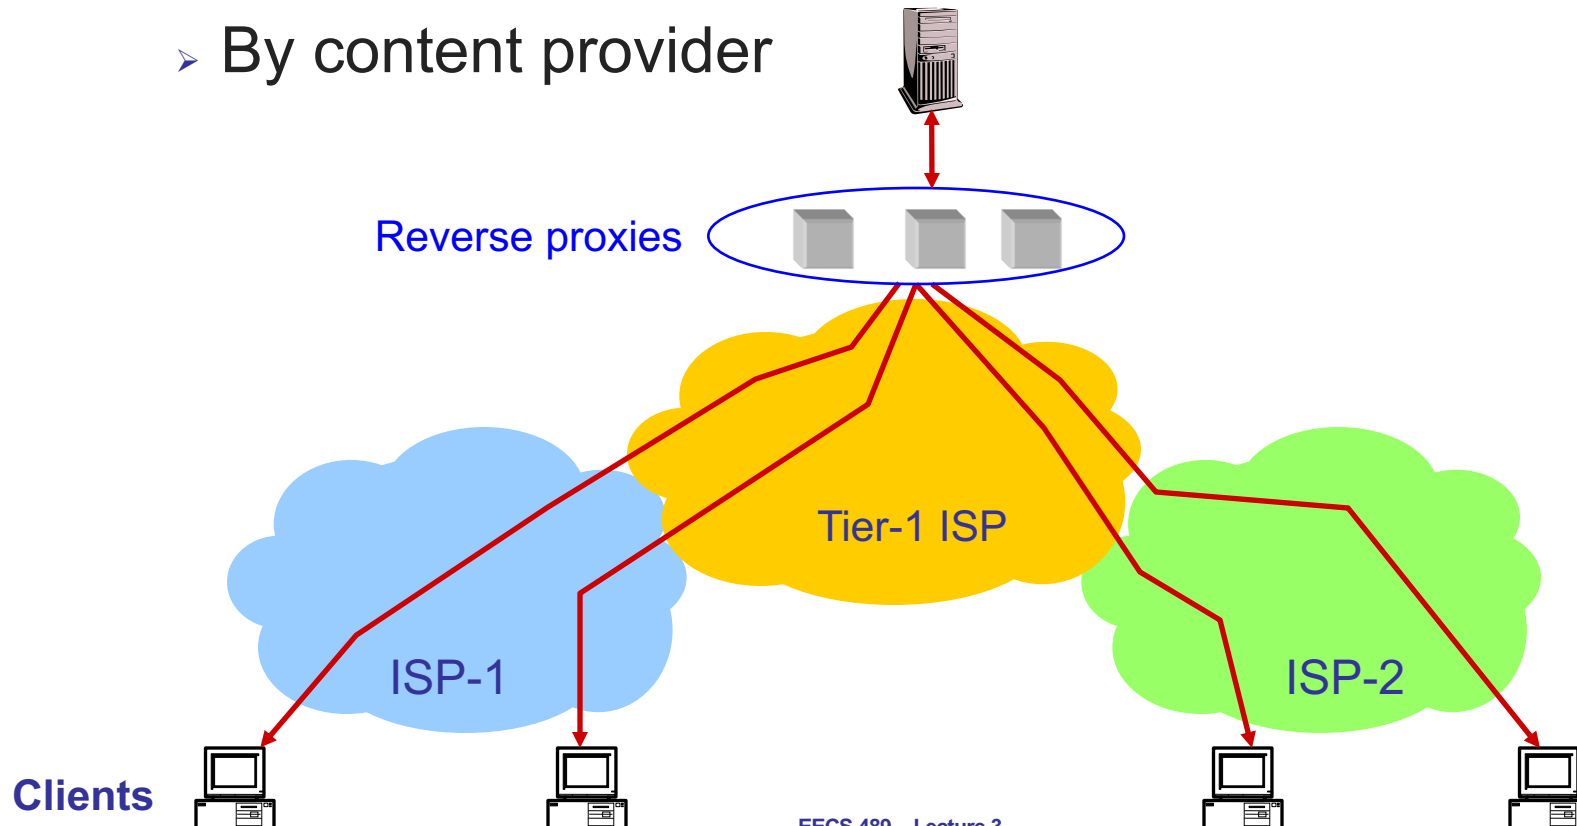
---

- Many clients transfer same information
  - Generate unnecessary server and network load
  - Clients experience unnecessary latency



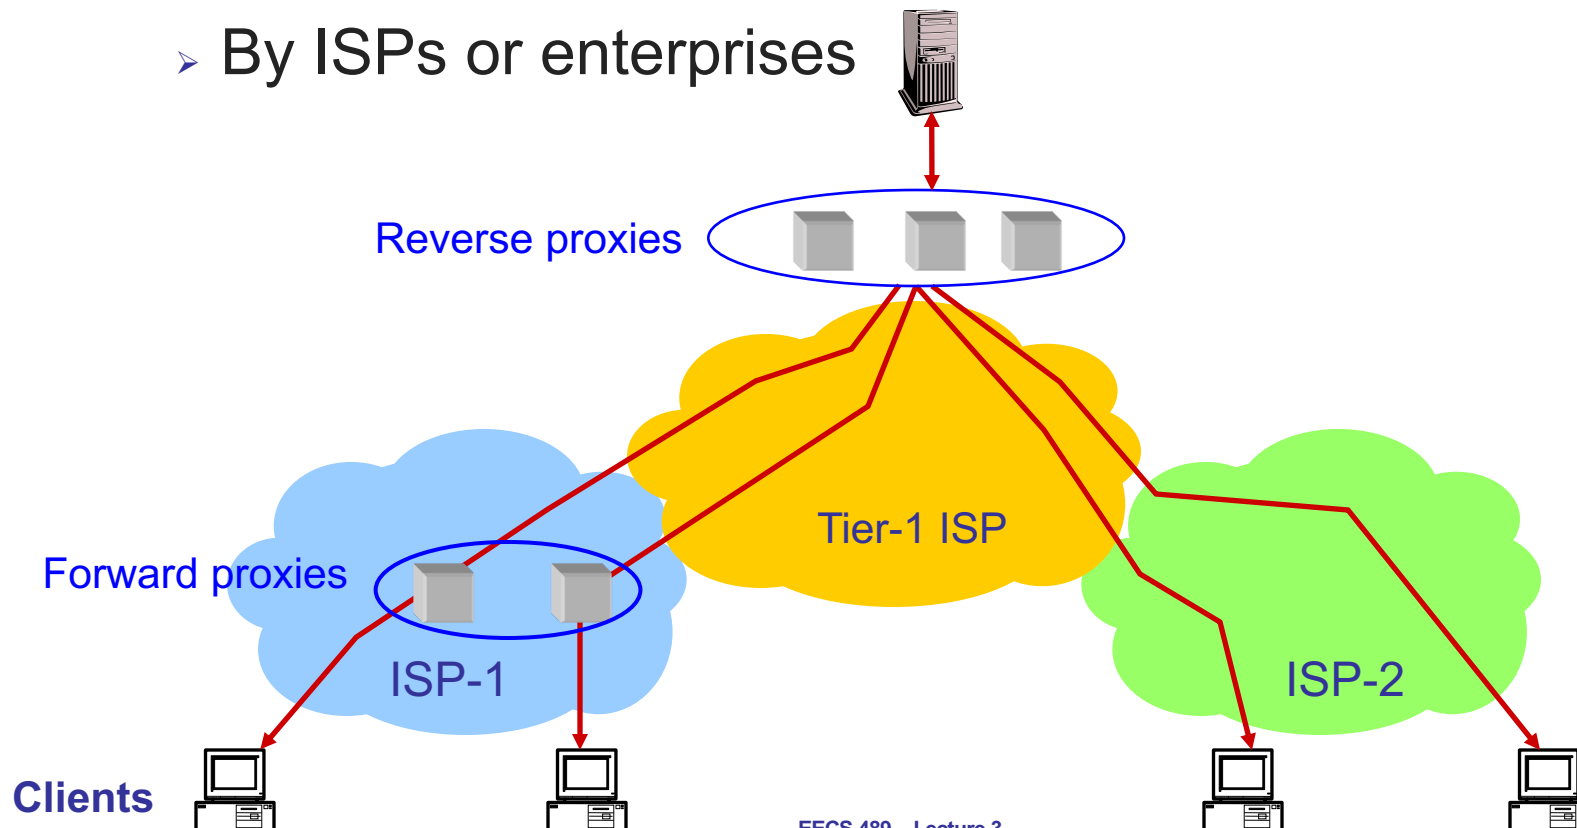
# Caching with Reverse Proxies

- Cache documents close to server
  - Decrease server load
  - By content provider



# Caching with Forward Proxies

- Cache documents close to clients
  - Reduce network traffic and decrease latency
  - By ISPs or enterprises



# Summary

---

- HTTP/1.1
  - Text-based protocol
  - Replaced by binary HTTP/2 protocol, which has been replaced by HTTP/3 in 2022
- Many ways to improve performance
  - Pipelining and batching
  - Caching in proxies and CDNs
  - Datacenters