

EECS 489

Computer Networks

Winter 2024

Mosharaf Chowdhury

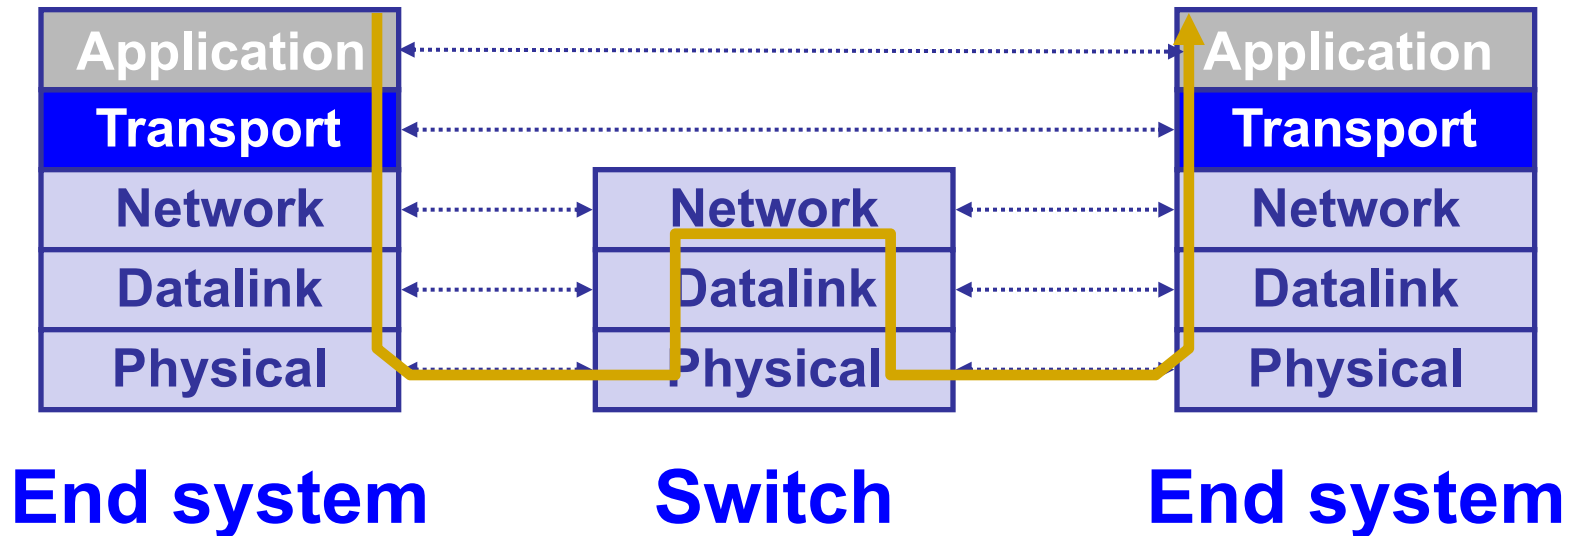
Material with thanks to Aditya Akella, Sugih Jamin, Philip Levis, Sylvia Ratnasamy, Peter Steenkiste, and many other colleagues.

Agenda

- Transport layer basics
- UDP
- Designing a reliable transport protocol

Transport layer

- Layer at **end hosts**, between the application and network layer



Why a transport layer?

- IP addresses capture hosts, but end-to-end communication happens between applications
 - Need a way to decide which packets go to which applications (multiplexing/demultiplexing)
- IP provides a weak service model (best-effort)
 - Packets can be corrupted, delayed, dropped, reordered, duplicated
 - No guidance on how much traffic to send and when
 - Dealing with this is tedious for application developers

Multiplexing & demultiplexing

- Multiplexing (Mux)
 - Gather and combining data chunks at the source from different applications and delivering to the network layer
- Demultiplexing (Demux)
 - Delivering correct data to corresponding sockets from a multiplexed stream

Role of the transport layer

- Communication between processes
 - Mux and demux from/to application processes
 - Implemented using *ports*

Role of the transport layer

- Communication between processes
- Provide common end-to-end services for app layer [optional]
 - Reliable, in-order data delivery
 - Well-paced data delivery
 - » Too fast may overwhelm the network
 - » Too slow is not efficient

Role of the transport layer

- Communication between processes
- Provide common end-to-end services for app layer [optional]
- TCP and UDP are the common transport protocols
 - Also SCTP, MPTCP, SST, RDP, DCCP, ...

Role of the transport layer

- Communication between processes
- Provide common end-to-end services for app layer [optional]
- TCP and UDP are the common transport protocols
- **UDP is a minimalist transport protocol**
 - Only provides mux/demux capabilities

Role of the transport layer

- Communication between processes
- Provide common end-to-end services for app layer [optional]
- TCP and UDP are the common transport protocols
- UDP is a minimalist transport protocol
- TCP offers a reliable, in-order, byte stream abstraction
 - With congestion control, but w/o performance guarantees (delay, b/w, etc.)

Applications and sockets

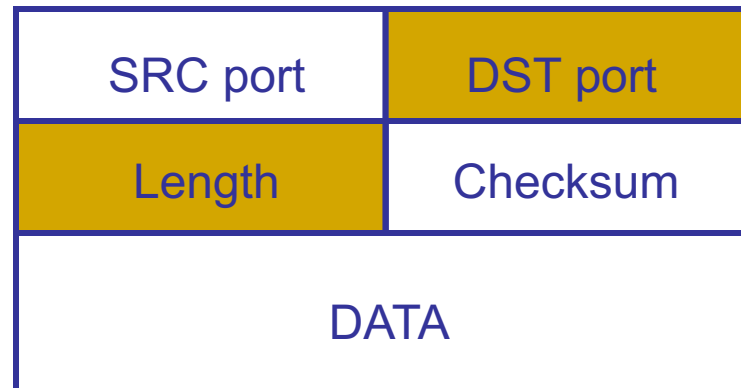
- **Socket**: software abstraction for an application process to exchange network messages with the (transport layer in the) operating system
- Two important types of sockets
 - UDP socket: TYPE is SOCK_DGRAM
 - TCP socket: TYPE is SOCK_STREAM

Ports

- 16-bit numbers that help distinguishing apps
 - Packets carry src/dst port no in transport header
 - Well-known (0-1023) and ephemeral ports
- OS stores mapping between sockets and ports
 - Port in packets and sockets in OS
 - For UDP ports (SOCK_DGRAM)
 - » OS stores (local port, local IP address) \leftrightarrow socket
 - For TCP ports (SOCK_STREAM)
 - » OS stores (local port, local IP, remote port, remote IP) \leftrightarrow socket

UDP: User Datagram Protocol

- Lightweight communication between processes
 - Avoid overhead and delays of order & reliability
- UDP described in RFC 768 – (1980!)
 - Destination IP address and port to support demultiplexing



UDP (cont'd)

- Optional error checking on the packet contents
 - (checksum field = 0 means “don’t verify checksum”)
- Source port is also optional
 - Useful to respond back to the sender in some cases

Why a transport layer?

- IP packets are addressed to a host but end-to-end communication is between application processes at hosts
 - Need a way to decide which packets go to which applications (mux/demux)
- IP provides a weak service model (best-effort)
 - Packets can be corrupted, delayed, dropped, reordered, duplicated
 - No guidance on how much traffic to send and when
 - Dealing with this is tedious for application developers

Reliable transport

- In a perfect world, reliable transport is easy

@Sender

- Send packets

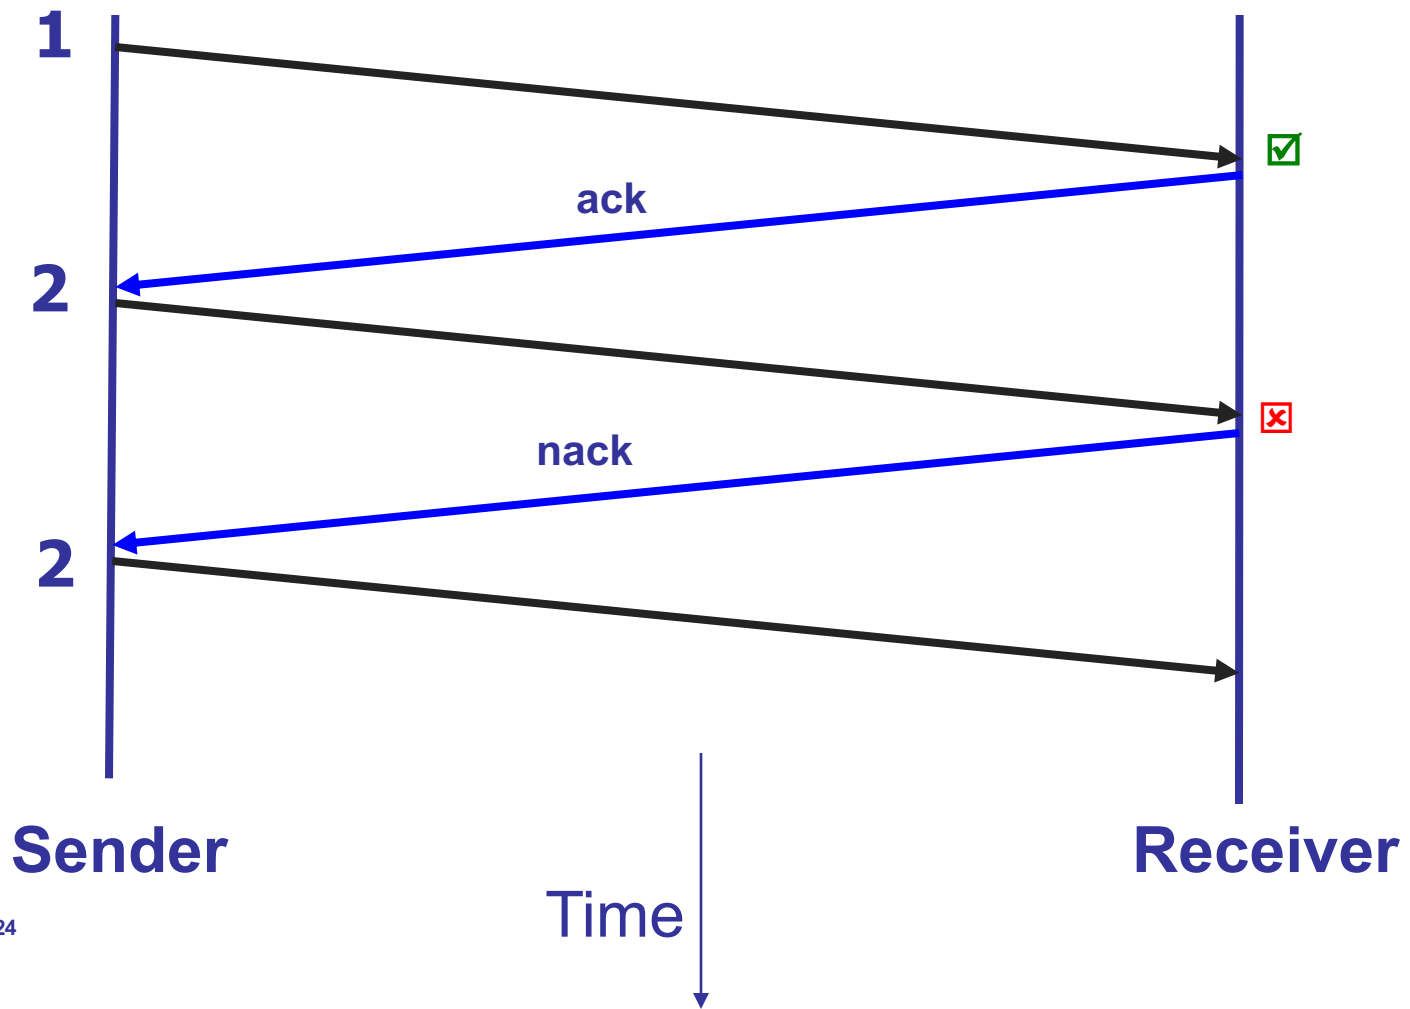
@Receiver

- Wait for packets

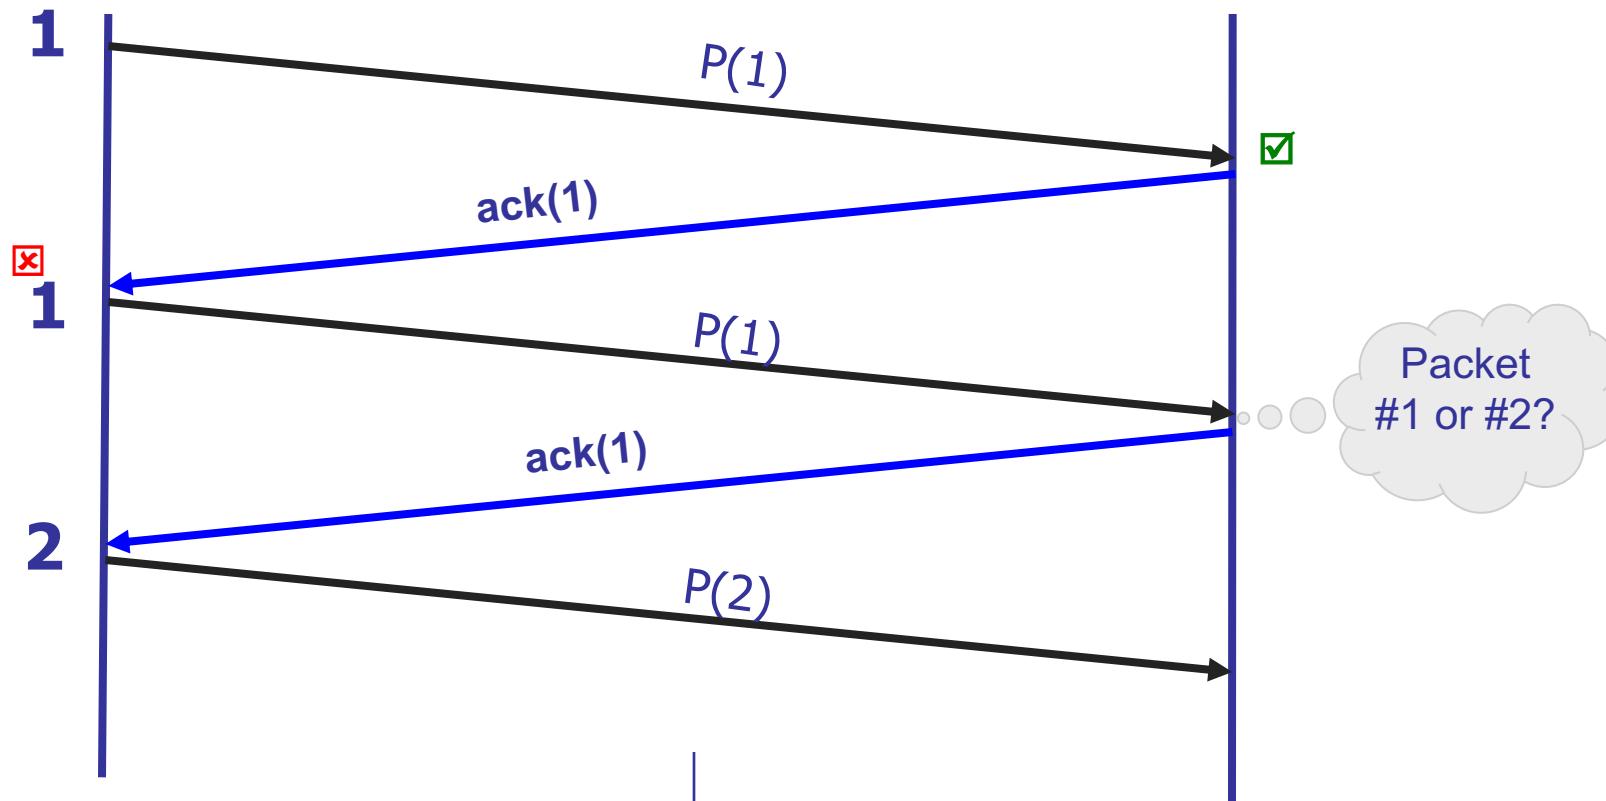
Reliable transport

- In a perfect world, reliable transport is easy
- All the bad things best-effort can do
 - A packet is corrupted (bit errors)
 - A packet is lost (*why?*)
 - A packet is delayed (*why?*)
 - Packets are reordered (*why?*)
 - A packet is duplicated (*why?*)

Dealing with packet corruption




Dealing with packet corruption



Data and ACK packets carry sequence numbers

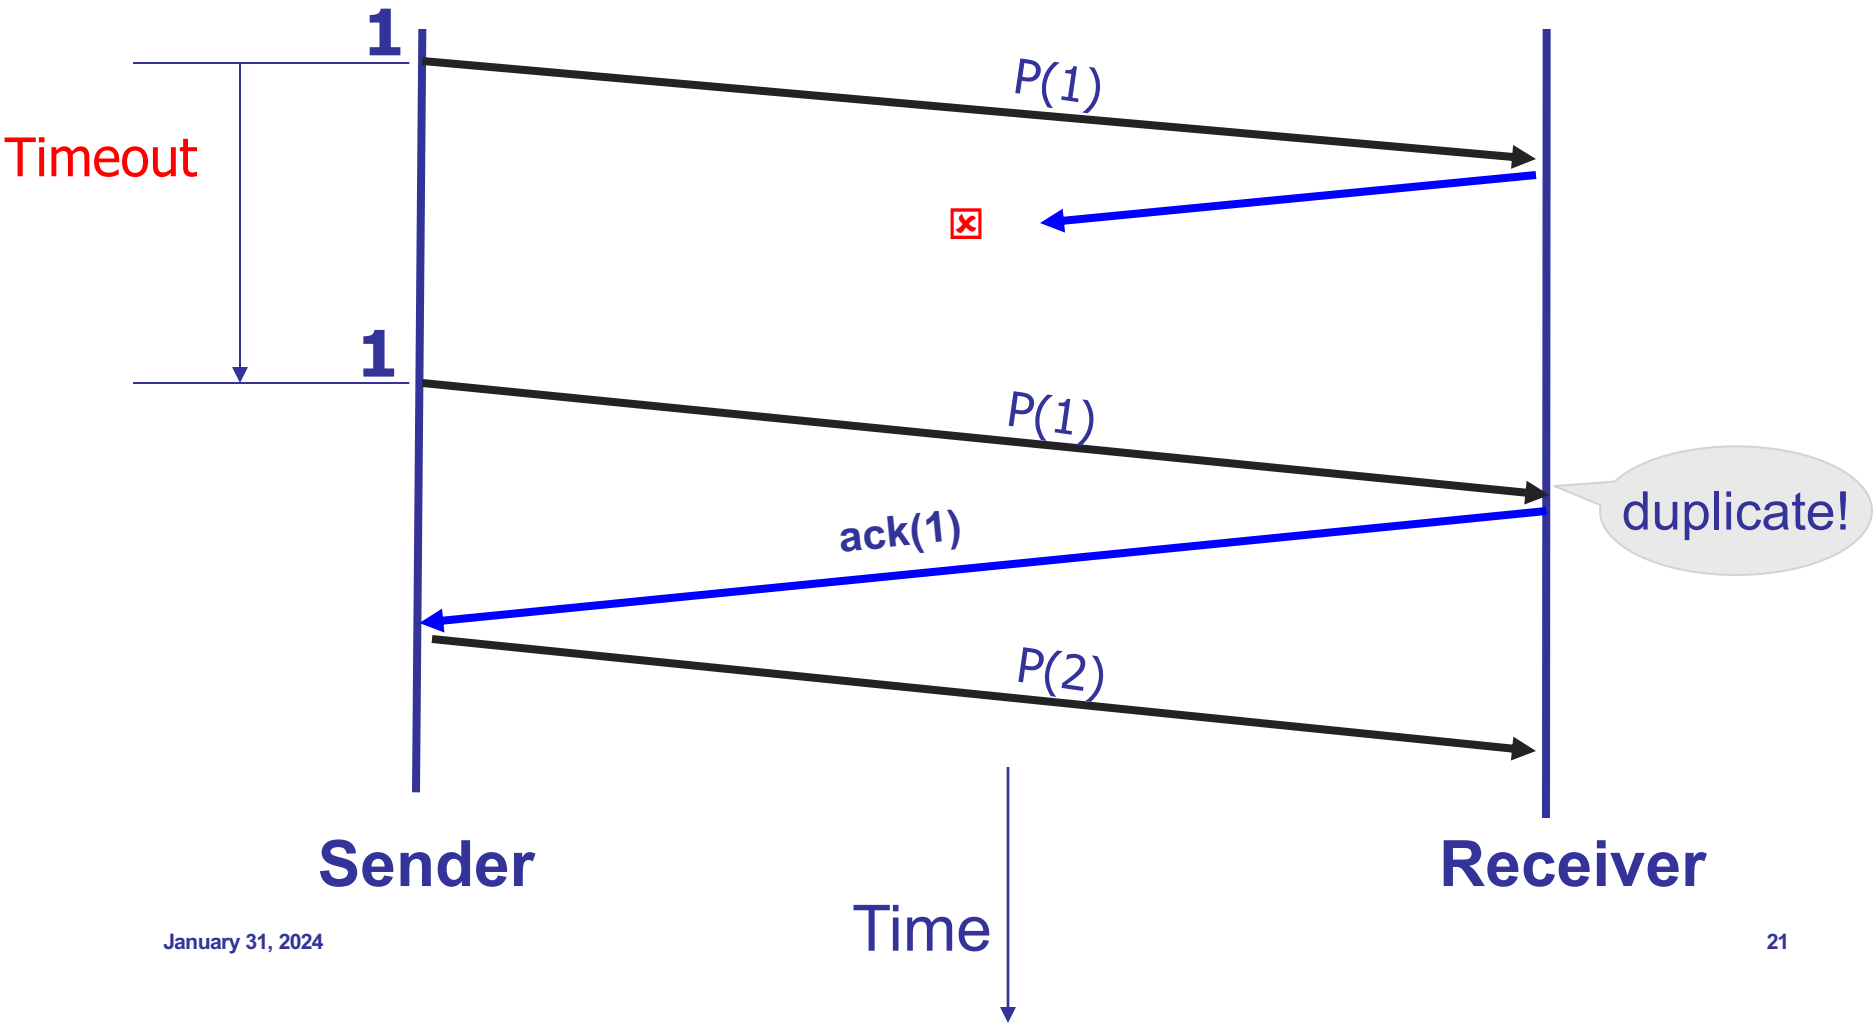
Dealing with packet loss



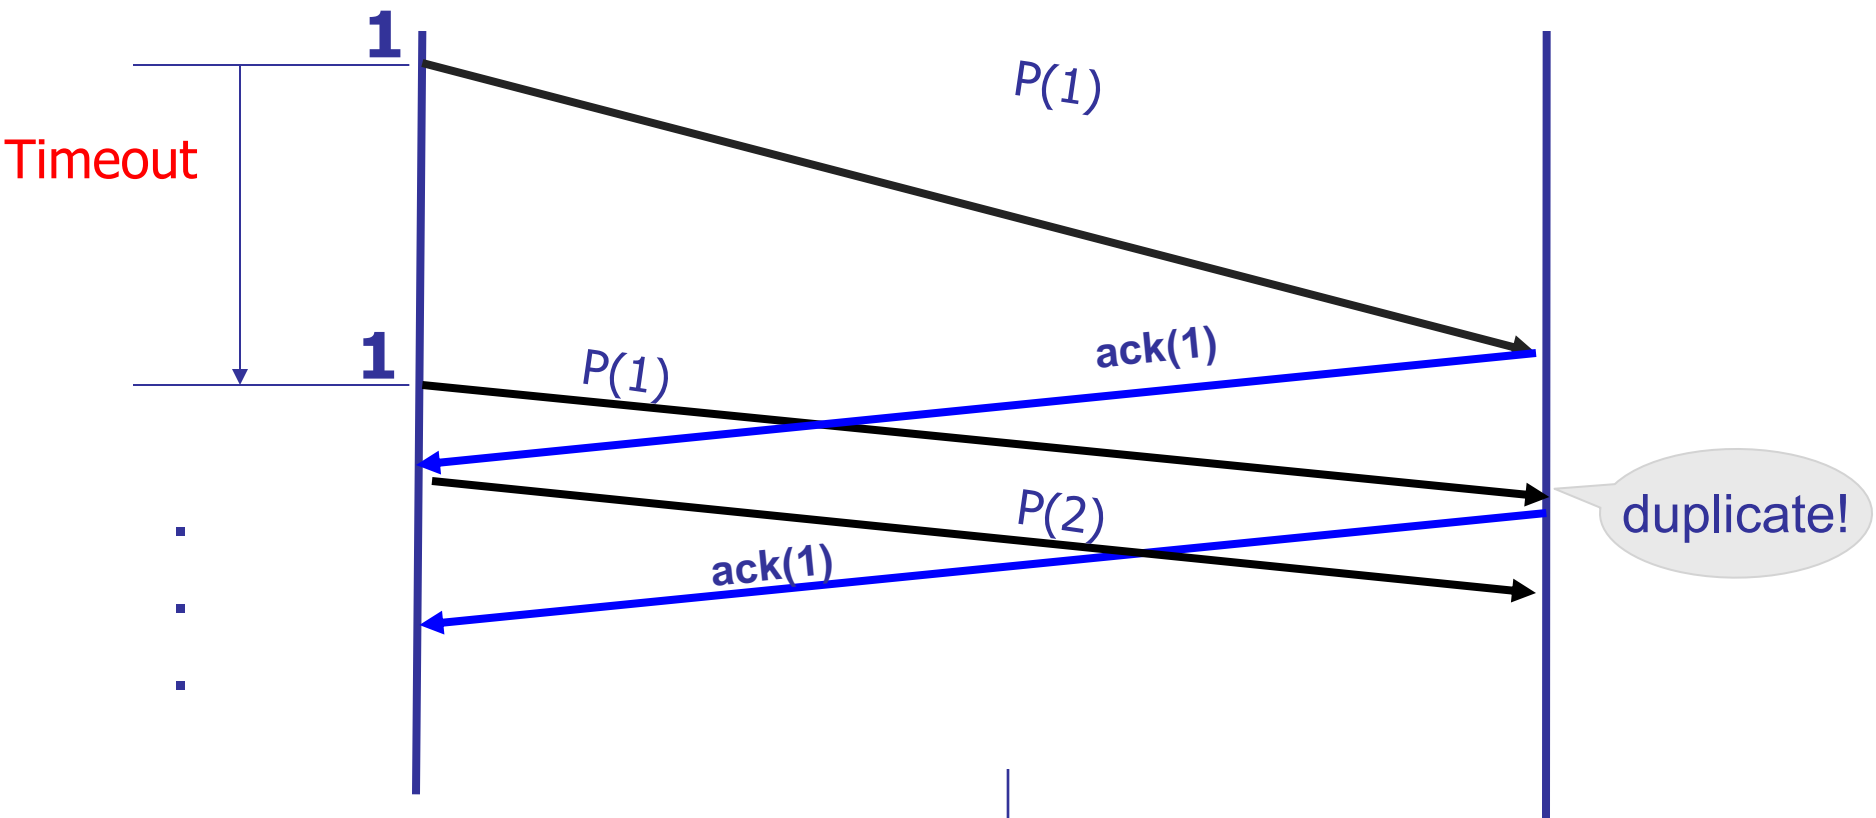
Timer-driven loss detection

Set timer when packet is sent; retransmit on timeout

Dealing with packet loss (of ack)



Dealing with delay



Timer-driven retransmission can lead to duplicates

Components of a solution

- Checksums (to detect bit errors)
- Timers (to detect loss)
- Acknowledgements (positive or negative)
- Sequence numbers (to deal with duplicates)

5-MINUTE BREAK!

DESIGNING A RELIABLE TRANSPORT

A Solution: “Stop and Wait”

@Sender

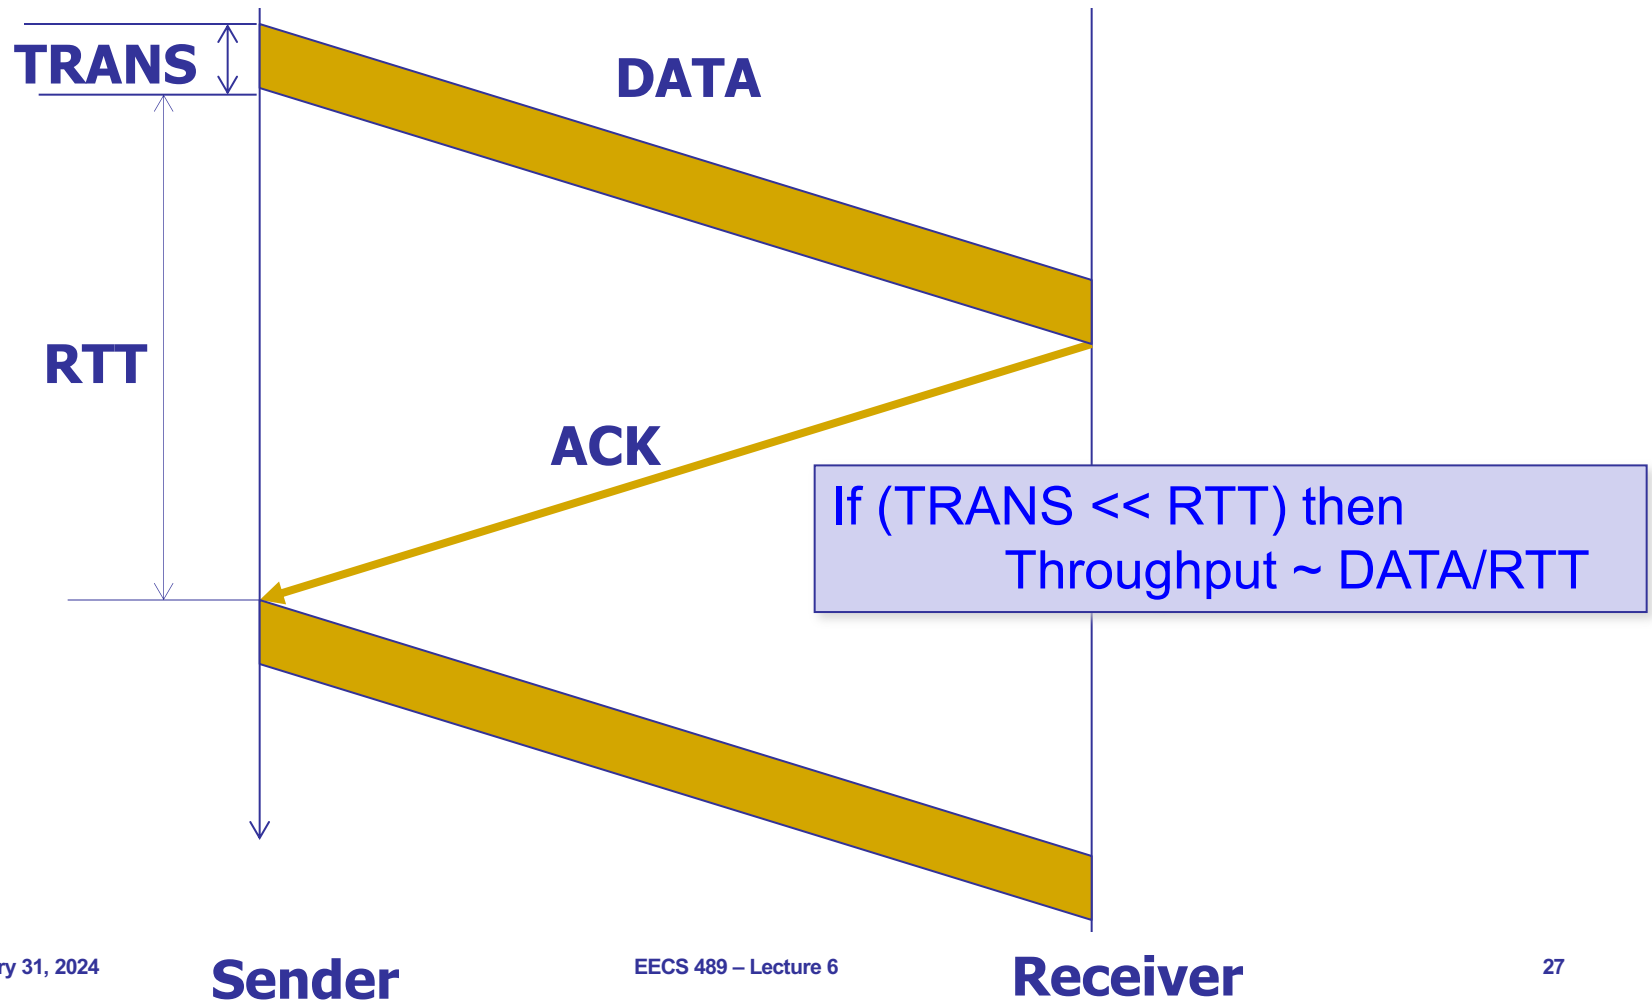
- Send packet(l); (re)set timer; wait for ack
- If (ACK)
 - l++; repeat
- If (NACK or TIMEOUT)
 - repeat

@Receiver

- Wait for packet
- If packet is OK, send ACK
- Else, send NACK or DISCARD
- Repeat

- A **correct** reliable transport protocol, but an **extremely inefficient** one

Stop & Wait is inefficient



Orders of magnitude

- Transmission time for 10Gbps link:
 - ~ microsecond for 1500 byte packet
- RTT:
 - 1,000 kilometers ~ $O(10)$ milliseconds

Three design decisions

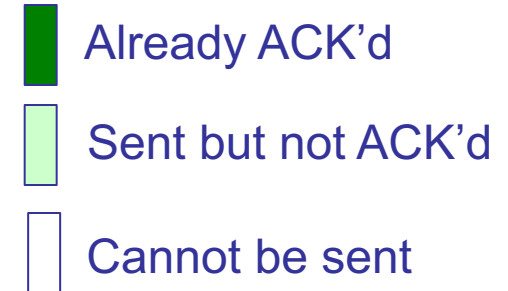
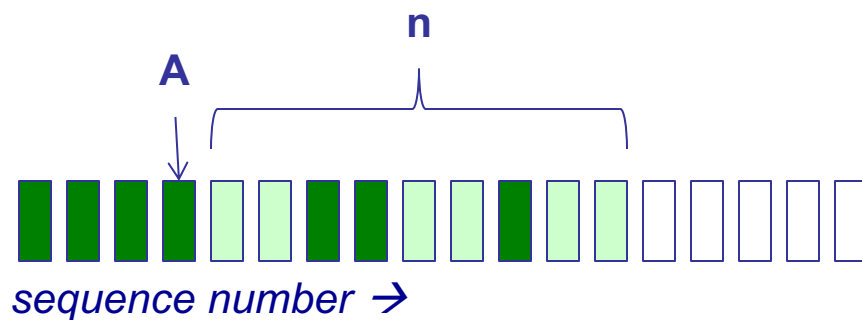
- Which **packets** can sender send?
- How does receiver ack packets?
- Which packets does sender resend?

Sliding window

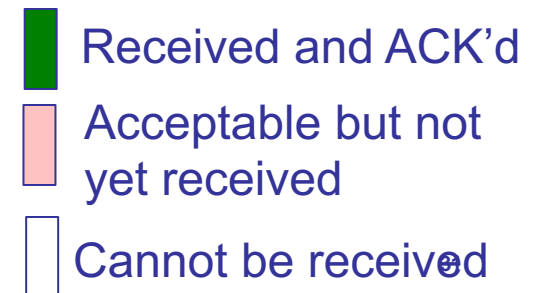
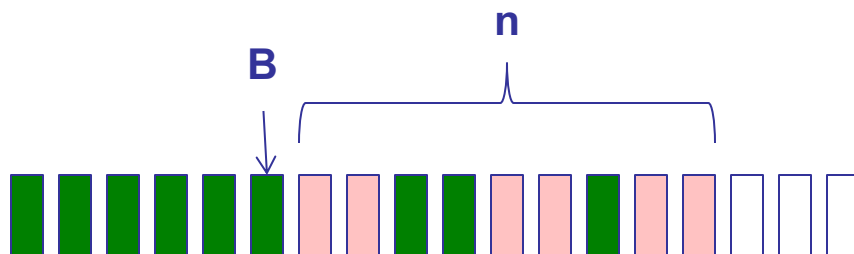
- Window = set of adjacent sequence numbers
 - The size of the set is the window size; assume window size is n
- General idea: send up to n packets at a time
 - Sender can send packets in its window
 - Receiver can accept packets in its window
 - Window of acceptable packets “slides” on successful reception/acknowledgement
 - Window contains all packets that might still be in transit
- Sliding window often called “packets in flight”

Sliding window

- Let A be the **last ack'd packet of sender without gap**;
then window of sender = $\{A+1, A+2, \dots, A+n\}$



- Let B be the **last received packet without gap** by receiver, then window of receiver = $\{B+1, \dots, B+n\}$



Throughput of sliding window

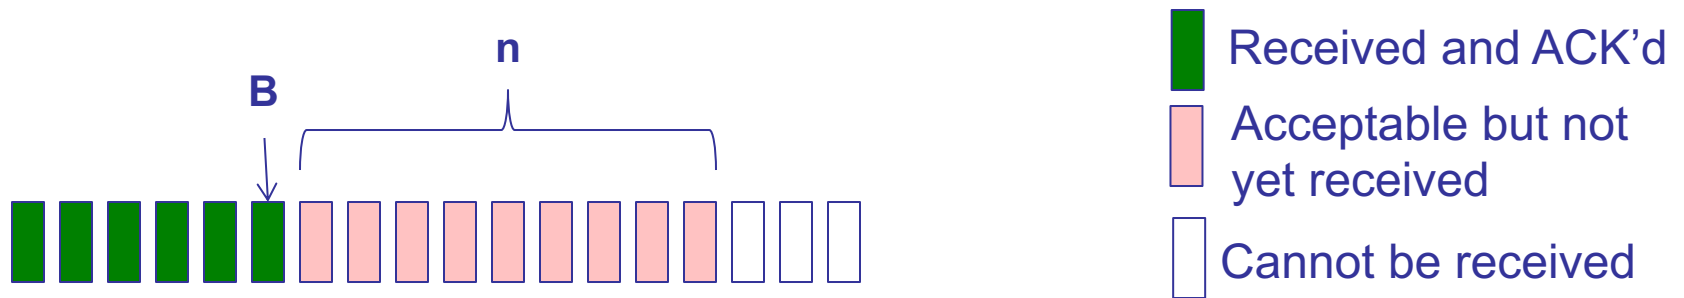
- If window size is n , then throughput is roughly
 - $\text{MIN}(n * \text{DATA} / \text{RTT}, \text{Link Bandwidth})$
- Compare to Stop and Wait: Data / RTT
- What happens when n gets too large?

Acknowledgements w/ sliding window

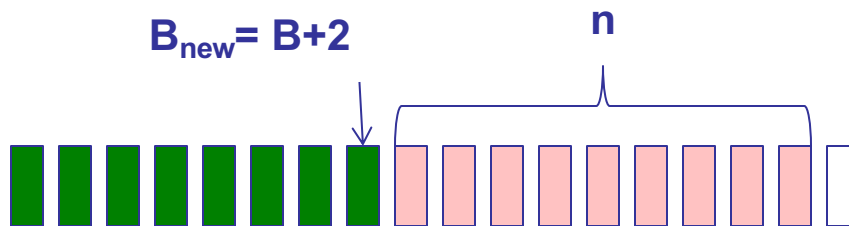
- Two common options
 - Cumulative ACKs: ACK carries next in-order sequence number that the receiver expects

Cumulative acknowledgements

- At receiver



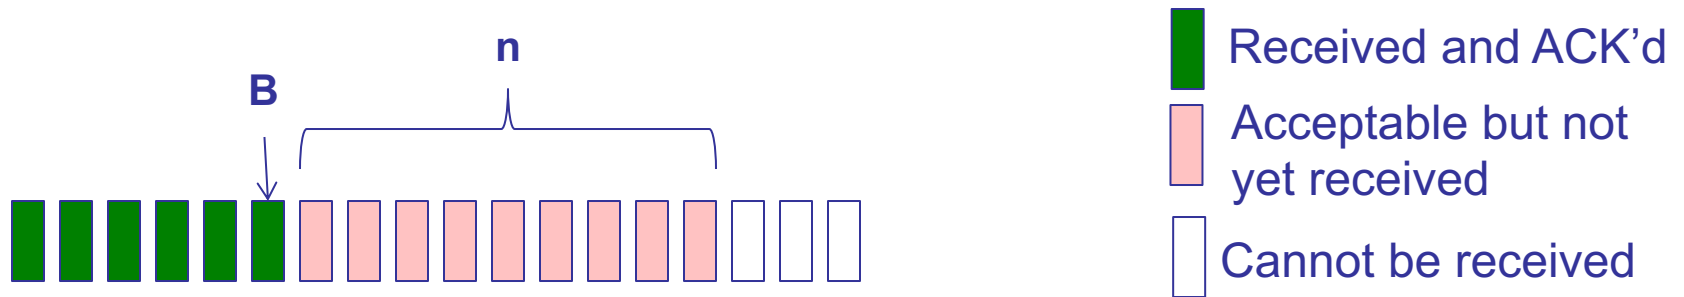
- After receiving B+1, B+2




- Receiver sends $\text{ACK}(B+3) = \text{ACK}(B_{\text{new}}+1)$

Cumulative acknowledgements (cont'd)

- At receiver



- After receiving B+4, B+5



- Receiver sends $ACK(B+1)$

Acknowledgements w/ sliding window

- Two common options
 - Cumulative ACKs: ACK carries next in-order sequence number the receiver expects
 - Selective ACKs: ACK individually acknowledges correctly received packets
- Selective ACKs offer more precise information but require more complicated book-keeping

Sliding window protocols

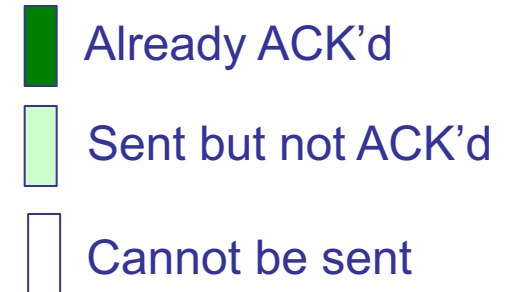
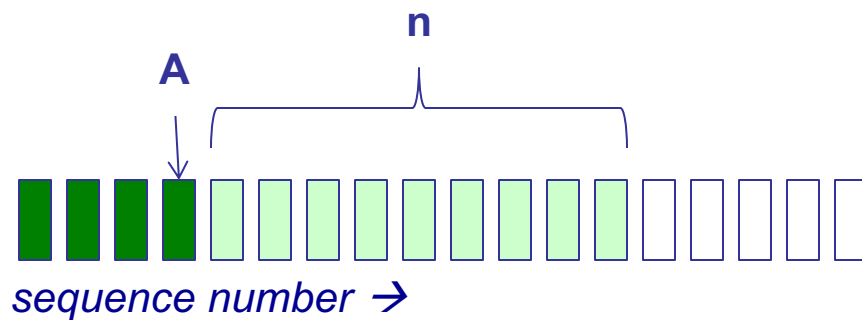
- Resending packets: two canonical approaches
 - Go-Back-N
 - Selective Repeat
- Many variants that differ in implementation details

Go-Back-N (GBN)

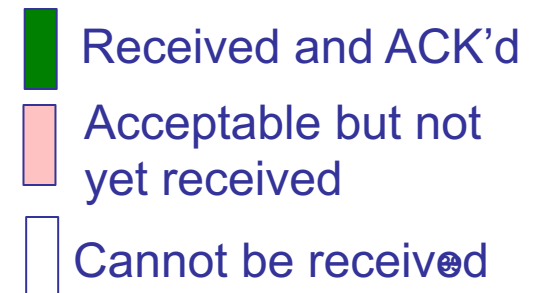
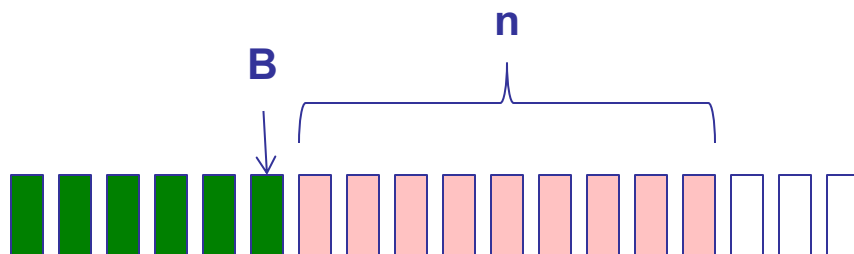
- Sender transmits up to n unacknowledged packets
- Receiver only accepts packets in order
 - Discards out-of-order packets (i.e., packets other than $B+1$)
- Receiver uses cumulative acknowledgements
 - i.e., sequence# in ACK = next expected in-order sequence#
- Sender sets timer for 1st outstanding ack ($A+1$)
- If timeout, retransmit $A+1, \dots, A+n$

Sliding window with GBN

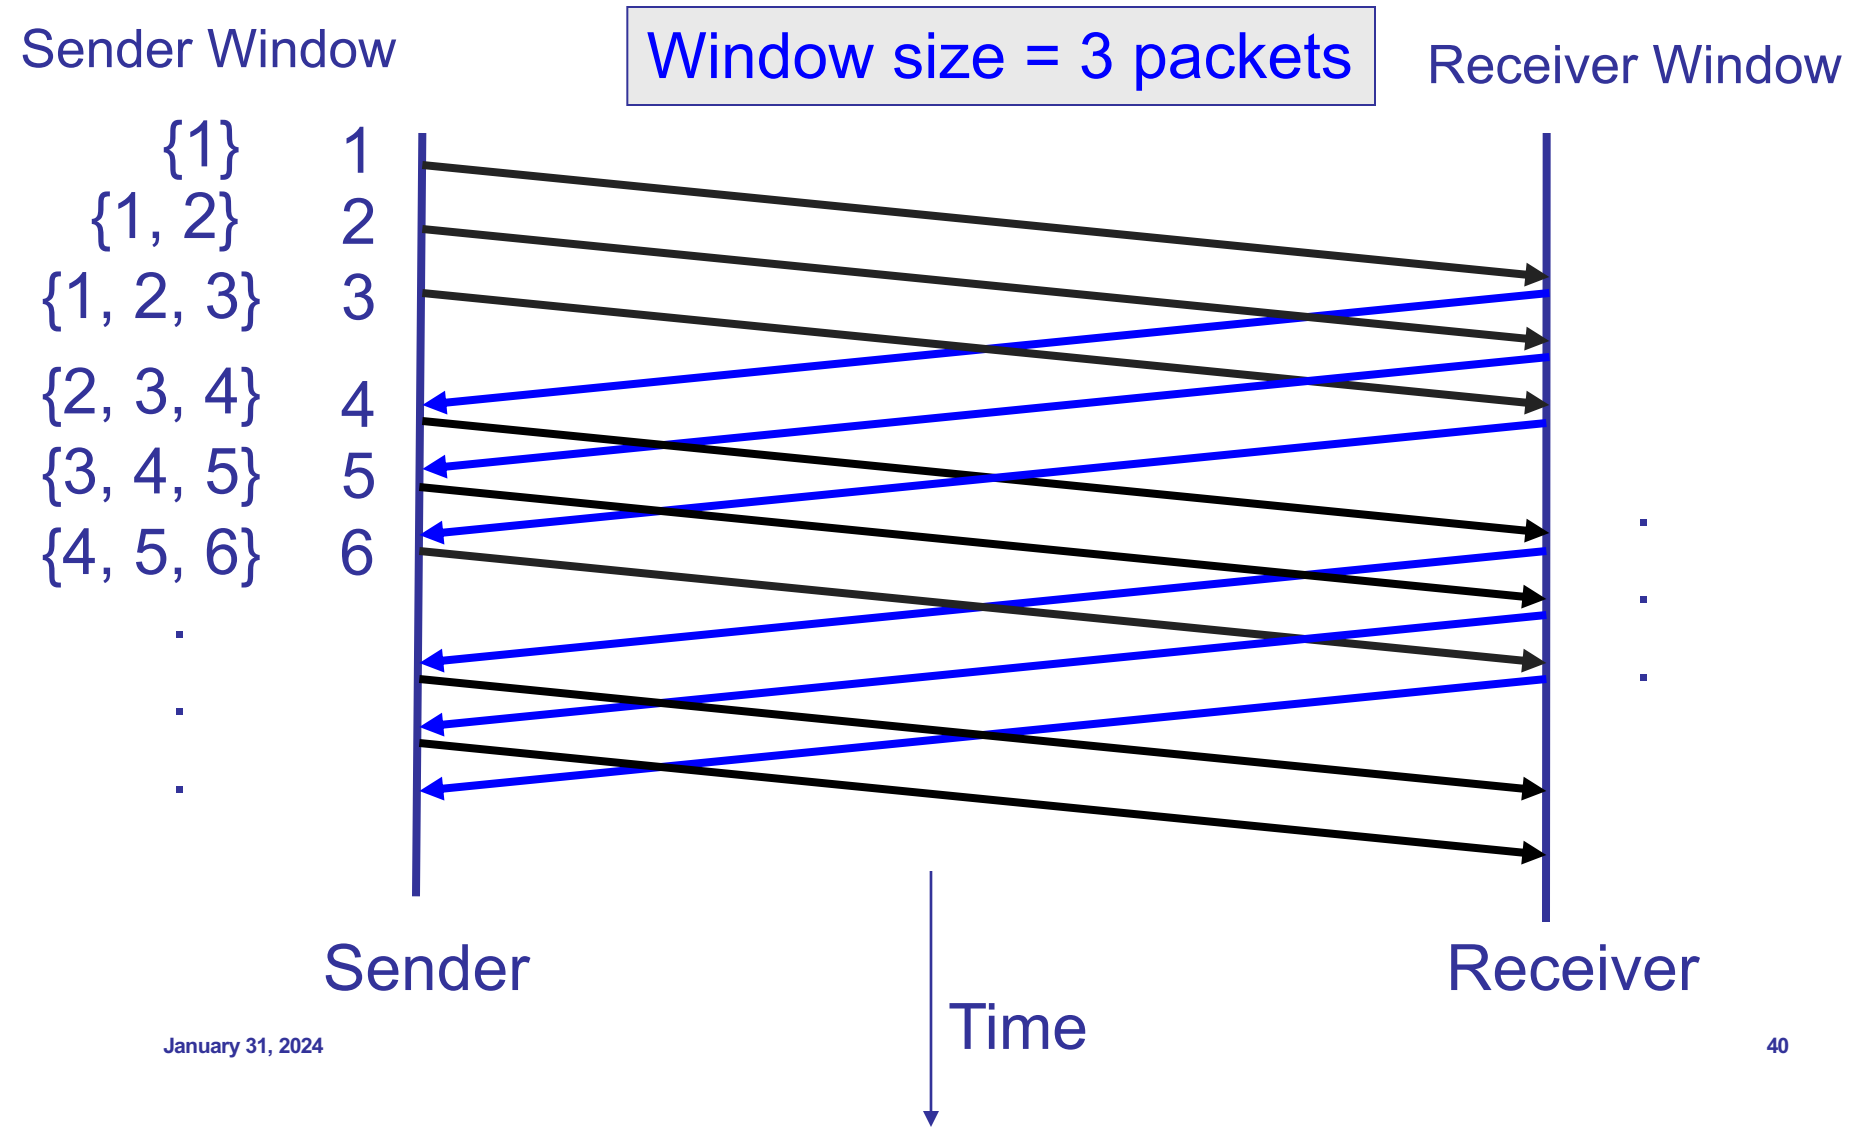
- Let A be the last ack'd packet of sender without gap;
then window of sender = $\{A+1, A+2, \dots, A+n\}$




- Let B be the last received packet without gap by receiver, then window of receiver = $\{B+1, \dots, B+n\}$



GBN example w/o errors



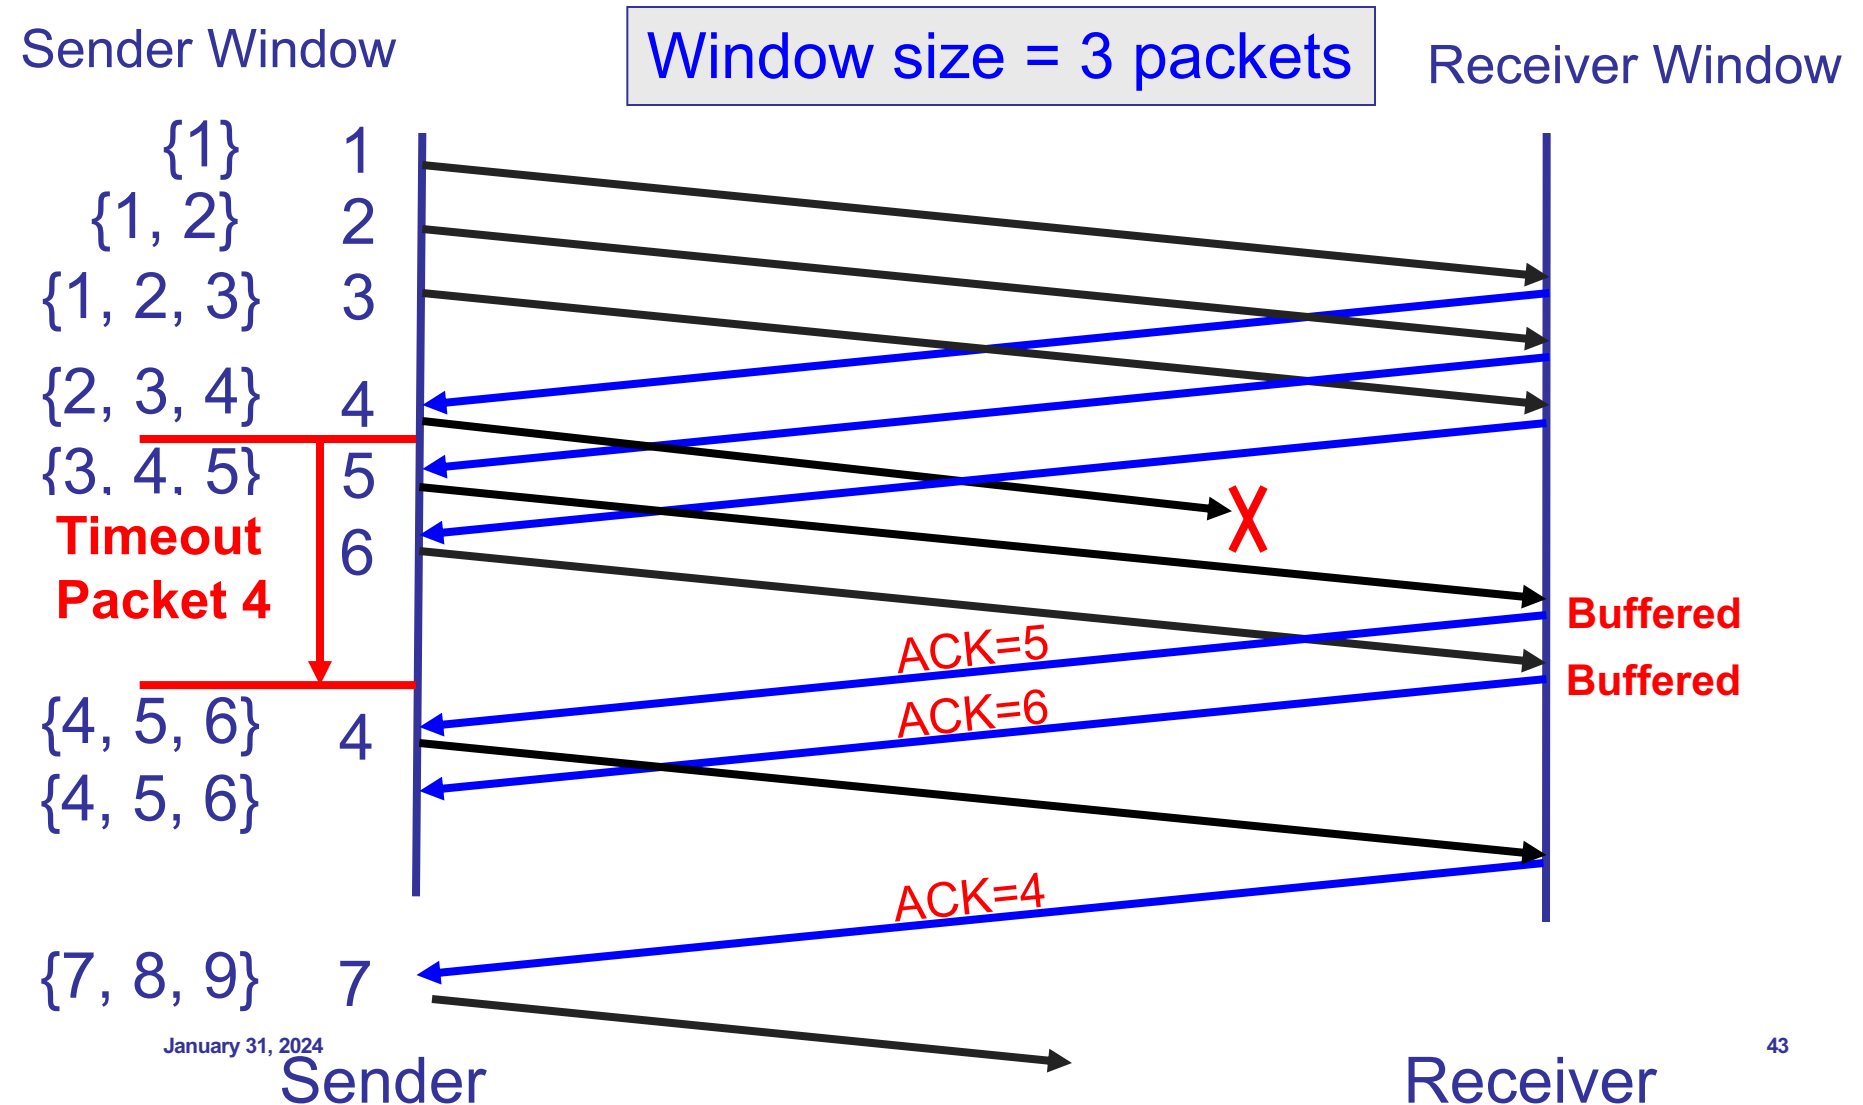
GBN example with errors



Selective Repeat (SR)

- Sender: transmit up to n unacknowledged packets
- Assume packet k is lost, $k+1$ is not
 - Receiver: indicates packet $k+1$ correctly received
 - Sender: retransmit only packet k on timeout
- Efficient in retransmissions but complex book-keeping
 - Need a timer per packet

SR example with errors



GBN vs. Selective Repeat

- When would GBN be better?
 - When error rate is low; wastes bandwidth otherwise
- When would SR be better?
 - When error rate is high; otherwise, too complex

Observations

- For a large-enough **window**, it is possible to fully utilize a link with sliding windows
- Sender has to **buffer** all unacknowledged packets, because they may require retransmission
- Receiver may be able to accept out-of-order packets, but only up to its **buffer** limits
- Implementation complexity depends on protocol details (GBN vs. SR)

Components of a solution

- Checksums (for error detection)
- Timers (for loss detection)
- Acknowledgments
 - Cumulative
 - Selective
- Sequence numbers (duplicates, windows)
- Sliding windows (for efficiency)
- Reliability protocols use the above to decide when and what to retransmit or acknowledge

Summary

- Transport layer allows applications to communicate with each other
- Provides unreliable and reliable mechanisms
- Possible to build reliable transport over unreliable medium

- Next lecture
 - TCP