

# **EECS 489**

# **Computer Networks**

**Winter 2024**

Mosharaf Chowdhury

*Material with thanks to Aditya Akella, Sugih Jamin, Xin Jin, Philip Levis, Sylvia Ratnasamy, Peter Steenkiste, and many other colleagues.*

# Agenda

---

- Software-defined networking
- Programmable networks

# The field of networking

---

- CS networking today is largely the study of the Internet
  - Perhaps the only history class many will take in CS

# Building an artifact, not a discipline

---

- Other fields in “systems”: OS, DB, etc.
  - Teach basic principles
  - Are easily managed
  - Continue to evolve
- Networking:
  - Teach big bag of protocols
  - Notoriously difficult to manage
  - Evolves very slowly
- Networks are much more primitive and less understood than other computer systems

# A tale of two planes

---

- **Data plane:** forwarding packets
  - Based on local forwarding state
- **Control plane:** computing that forwarding state
  - Involves coordination with rest of system

# Original goals for the control plane

---

- **Basic connectivity:** route packets to destination
  - Local state computed by routing protocols
  - Globally distributed algorithms
- **Inter-domain policy:** find policy-compliant paths
  - Done by globally distributed BGP
- What other goals are there in running a network?

# Extended roles of the control plane

---

- Performs various network management tasks
  - For example,
    - » Where to route?
    - » How much to route?
    - » At what rate to route?
    - » Should we route at all?
    - » ...

# Bottom line

---

- Many different control plane mechanisms
- Each designed from scratch for their intended goal
- Encompassing a wide variety of implementations
  - Distributed, manual, centralized,...
- None of them particularly well designed
- Network control plane is a complicated mess!



# “The Power of Abstraction”

---

- “Modularity based on abstraction is the way things get done”
  - Barbara Liskov
- Abstractions → Interfaces → Modularity

# Analogy: Mainframe to PC evolution

---

## Vertical integration, closed

- Specialized application
- Specialized operating system
- Specialized hardware

## Open interfaces

- Arbitrary applications
- Commodity operating systems
- Microprocessor

**We want the same for networking!**

# Many control plane mechanisms

---

- Variety of goals, no modularity
  - **Routing**: distributed routing algorithms
  - **Isolation**: ACLs, Firewalls,...
  - **Traffic engineering**: adjusting weights,...
- Control Plane: mechanism without abstraction
  - Too many mechanisms, not enough functionality

# Task: Compute forwarding state

---

- Consistent with low-level hardware/software
  - Which might depend on vendor
- Based on entire network topology
  - Because many control decisions depend on topology
- For all routers/switches in network
  - Every router/switch needs forwarding state

# Separate concerns with abstractions

---

- Be compatible with low-level hardware/software
  - Forwarding abstraction
- Make decisions based on entire network
  - Network state abstraction
- Compute configuration of each physical device
  - Specification abstraction

# #1: Forwarding abstraction

---

- Express intent independent of implementation
  - Don't want to deal with proprietary HW and SW
- Design details concern exact nature of:
  - Header matching
  - Allowed actions

# #2: Network state abstraction

---

- Abstraction: **global network view**
  - Annotated network graph provided through an API
- Creates a logically centralized view of the network (Network Operating System)
  - Runs on replicated servers in network (“controllers”)
- Information flows both ways
  - Information from routers/switches to form “view”
  - Configurations to routers/switches to control forwarding

# #3: Specification abstraction


---

- Control mechanism expresses desired behavior
  - Whether it be isolation, access control, or QoS
- It should not be responsible for implementing that behavior on physical network infrastructure
  - Requires configuring the forwarding tables in each switch
- **Abstract view** of network
  - Models only enough detail to specify goals
  - Will depend on task semantics



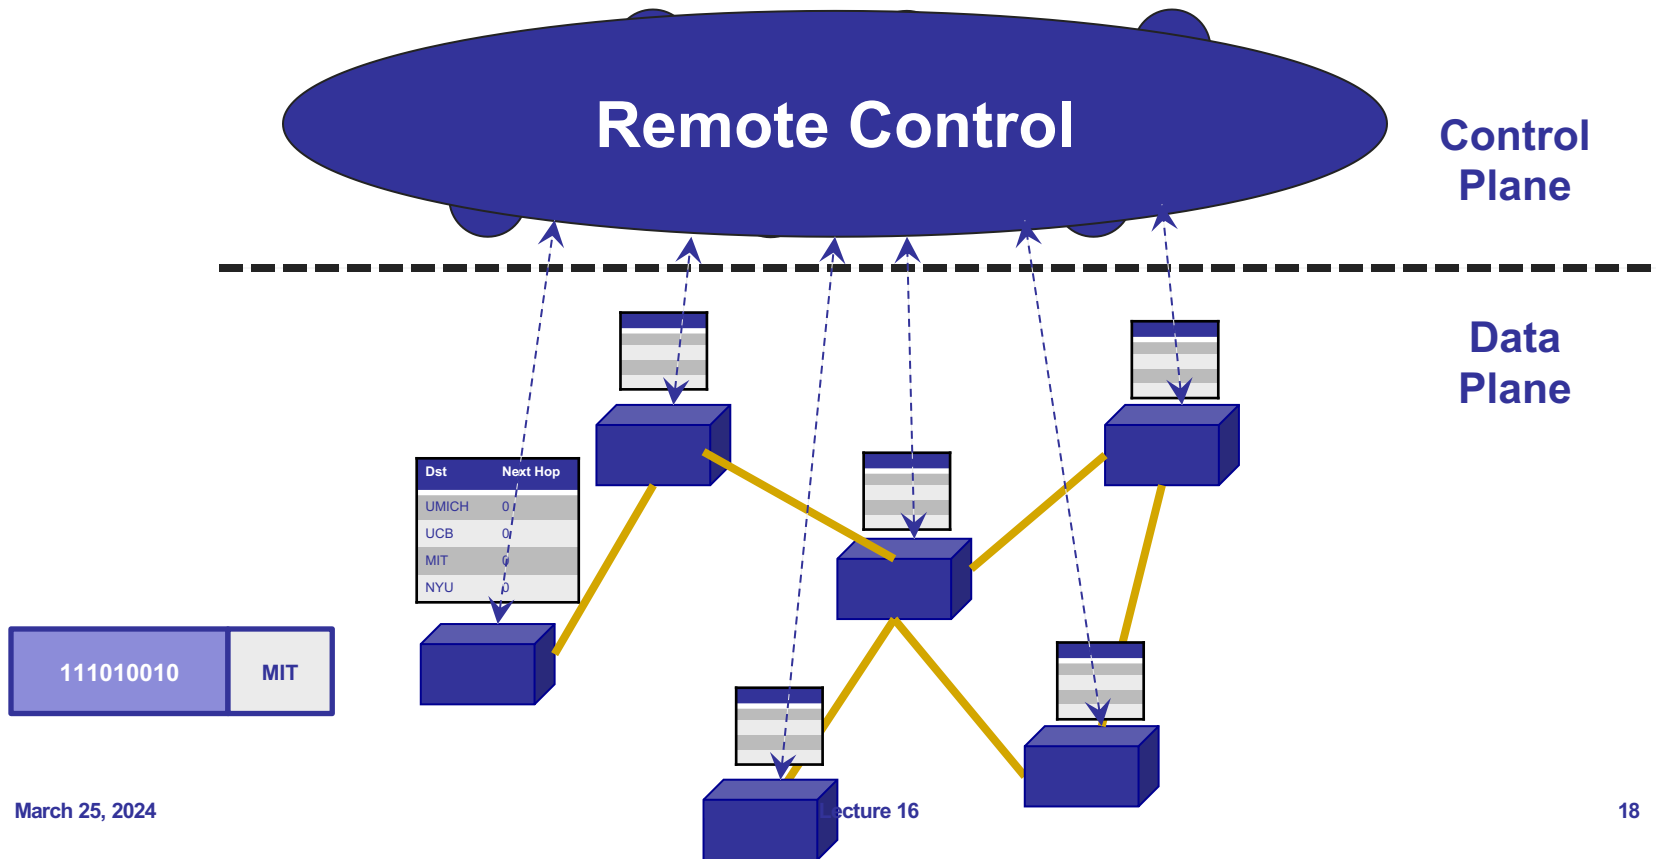
# Traditional fully decentralized control plane

- Individual routing algorithm components in every router interact in the control plane



# Logically centralized control plane

- A distinct (typically remote) controller interacts with local control agents (CAs)



# Each goal is an app via specification abstraction

---

- What if an operator wants X?
- What if a customer wants to do weighted traffic splitting?
- ...
- **There is an app for it!**
  - Write your own routing protocol, load balancing algorithm, access control policies

# Reason about each app via network state abstraction

---

- Now that the network is not distributed anymore and is a simple graph, we can **verify** whatever the correctness of whatever we specified

# Logically centralized control plane

---

- A distinct (typically remote) controller interacts with local control agents (CAs)
- Each router contains a **flow table**
- Each entry of the flow table defines a **match-action** rule
- Entries of the flow table is computed and distributed by the (logically) centralized controller

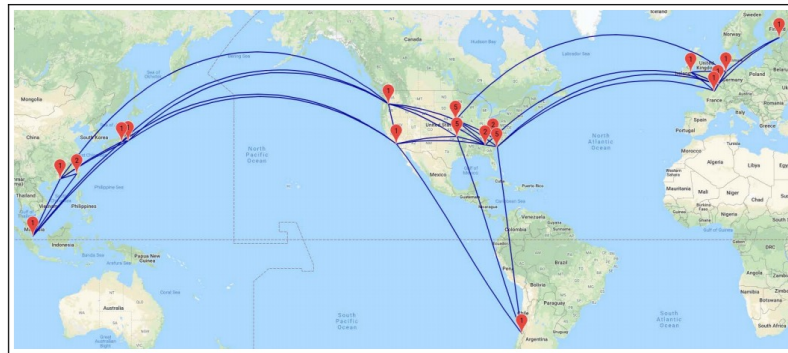
# SDN: Many challenges remain

---

- Hardening the control plane: dependable, reliable, performance-scalable, secure distributed system
  - Robustness to failures: leverage strong theory of reliable distributed system for control plane
  - Dependability, security: “baked in” from day one?
- Networks, protocols meeting mission-specific requirements
  - E.g., real-time, ultra-reliable, ultra-secure
- Internet-scaling

# Some progress in the wide-area network (WAN)

- Google and Microsoft use SDN to manage traffic between datacenters
- One centralized controller to rule the entire world (well, their world)



**Figure 1: B4 global topology.** Each marker indicates a site or multiple sites located in close geographical proximity. B4 consists of 33 sites as of January, 2018.

---

**5-MINUTE BREAK!**



# A tale of two planes

---

- **Data plane:** forwarding packets
  - Based on local forwarding state
- **Control plane:** computing that forwarding state
  - Involves coordination with rest of system

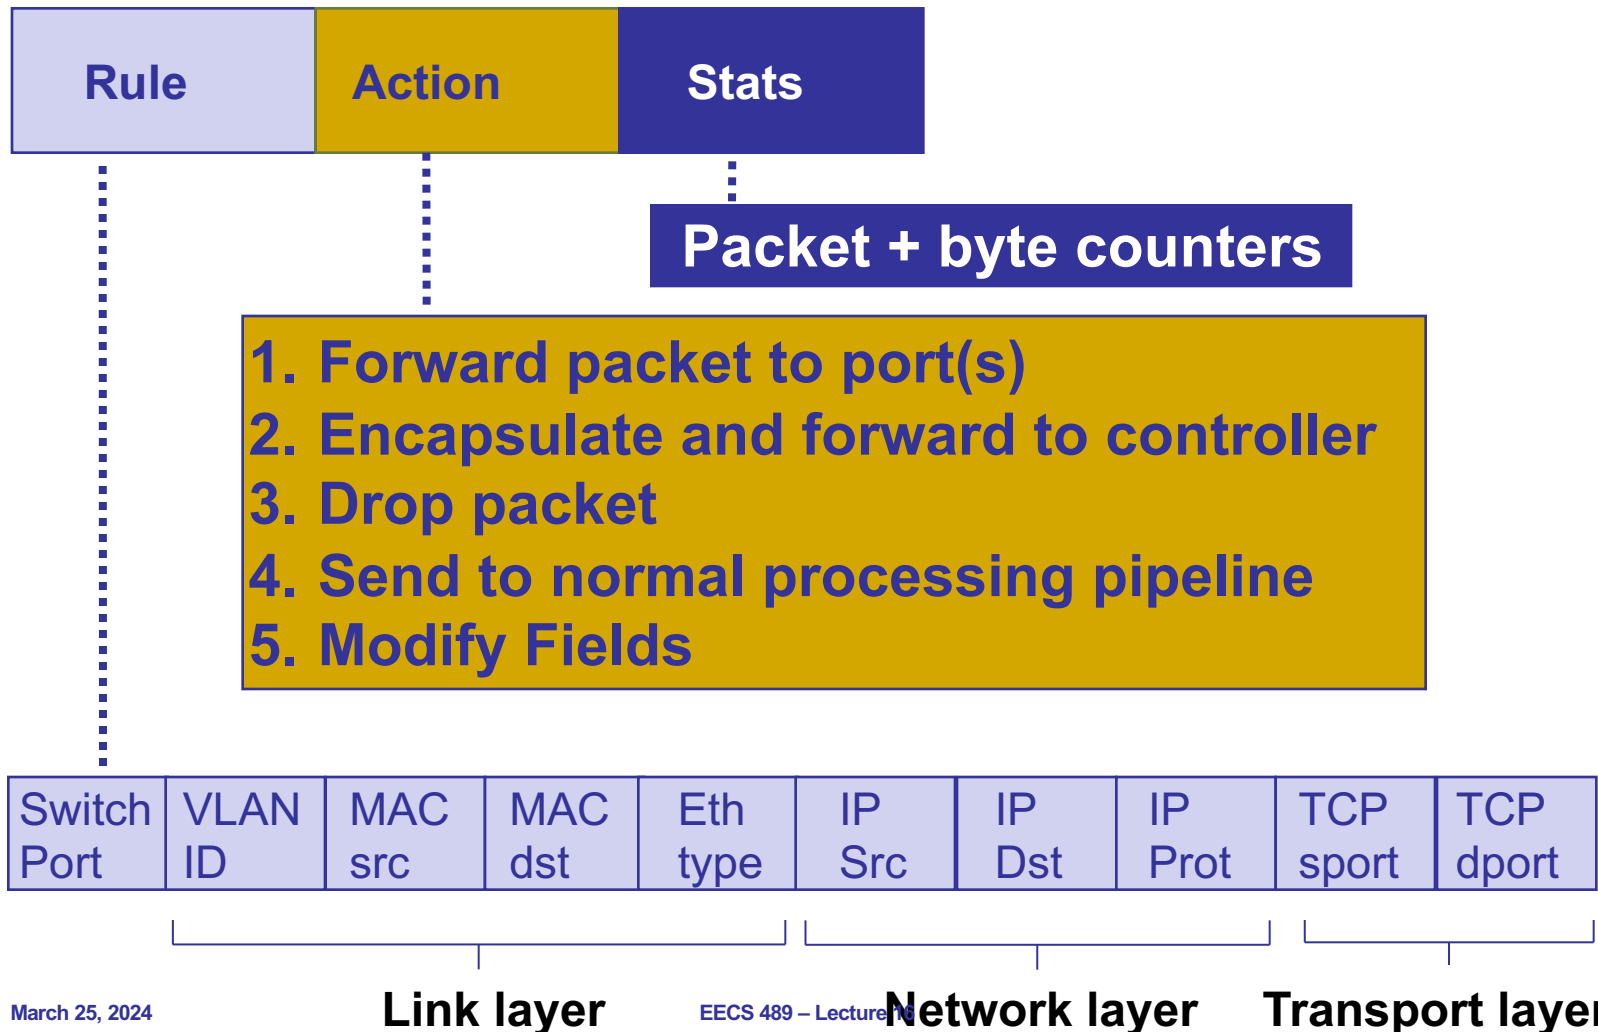
# OpenFlow data plane abstraction

---

- Flow is defined by header fields
- Generalized forwarding: simple packet-handling rules
  - **Pattern**: match values in packet header fields
  - **Actions**: for matched packet: drop, forward, modify, matched packet or send matched packet to controller
  - **Priority**: disambiguate overlapping patterns
  - **Counters**: #bytes and #packets

1. src=1.2.\*.\*, dest=3.4.5.\* → drop
2. src = \*.\*.\*.\*, dest=3.4.\*.\* → forward(2)
3. src=10.1.2.3, dest=\*.\*.\*.\* → send to controller

# OpenFlow: Flow table entries



# Forwarding abstraction

---

**Match + Action:** unifies different kinds of devices

- Router

- Match: longest destination IP prefix
- Action: forward out a link

- Switch

- Match: destination MAC address
- Action: forward or flood

- Firewall


- Match: IP addresses and TCP/UDP port numbers
- Action: permit or deny

- NAT

- Match: IP address and port
- Action: rewrite address and port

# OpenFlow example

match	action
IP Src = 10.3.*.* IP Dst = 10.2.*.*	forward(3)




**Example:** datagrams from hosts h5 and h6 should be sent to h3 or h4, via s1 and from there to s2

match	action
ingress port = 1 IP Src = 10.3.*.* IP Dst = 10.2.*.*	forward(4)

match	action
ingress port = 2 IP Dst = 10.2.0.3	forward(3)
ingress port = 2 IP Dst = 10.2.0.4	forward(4)

# OpenFlow protocol

## OpenFlow Controller



- Operates between controller, switch
- TCP used to exchange messages
  - Optional encryption
- Three classes of OpenFlow messages:
  - Controller-to-switch
  - Asynchronous (switch to controller)
  - Symmetric (misc.)

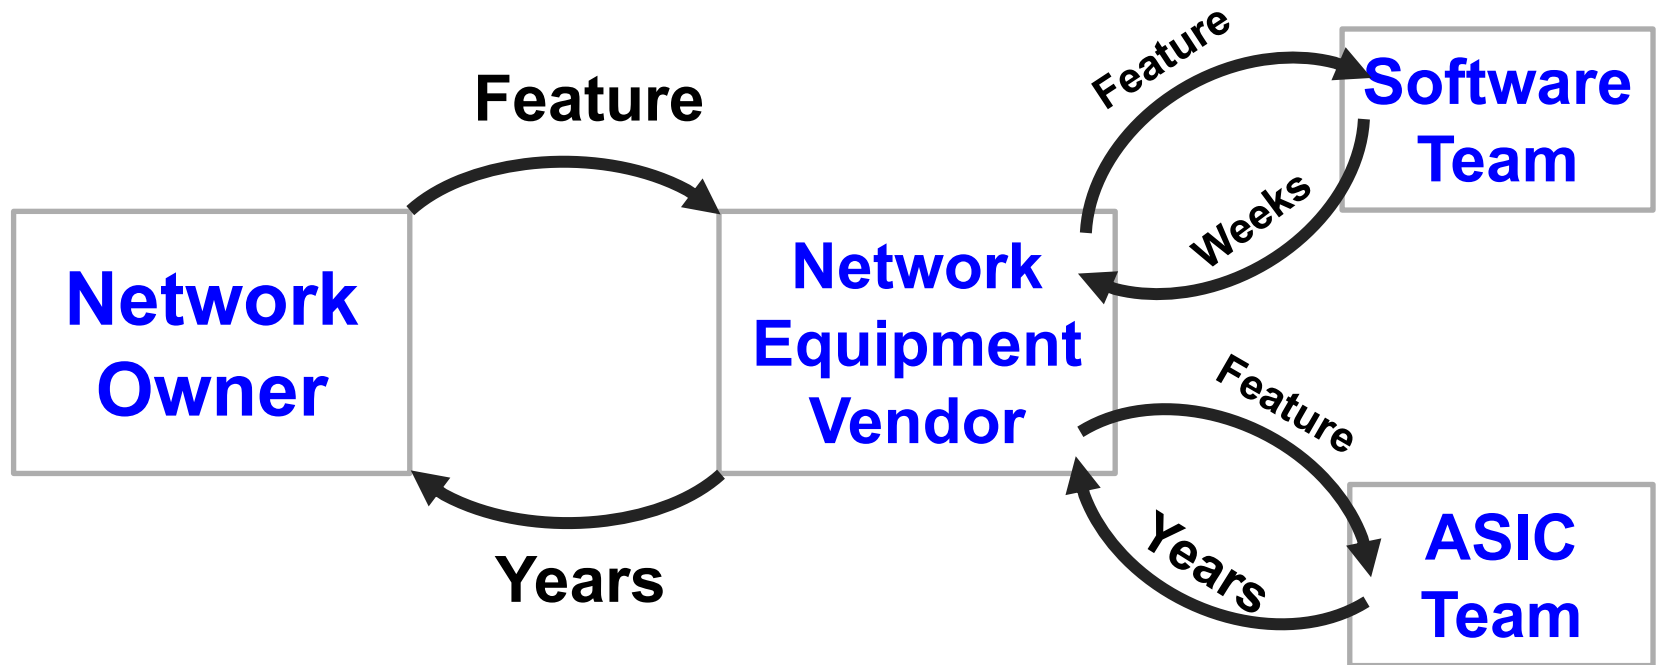
# Fixed-function data plane

---

- Traditional switches are fixed-function
  - They can do whatever they can do at birth, but they cannot change!
  - Bottom-up design
- Even OpenFlow was designed to be a fixed protocol
  - With a fixed table format
  - Capable of doing limited things

# Takes forever to get a feature

---





# Programmable data plane

---

- What if we could tell switches exactly what we want?
  - What table to keep?
  - What rules to use?
  - What data to keep track of?
  - ...

# Top-down approach

---

- **Precisely specify** what you want to do and how you want a packet to be processed

```
table int_table {  
    reads {  
        ip.protocol  
    }  
    actions {  
        export_queue_latency  
    }  
}
```

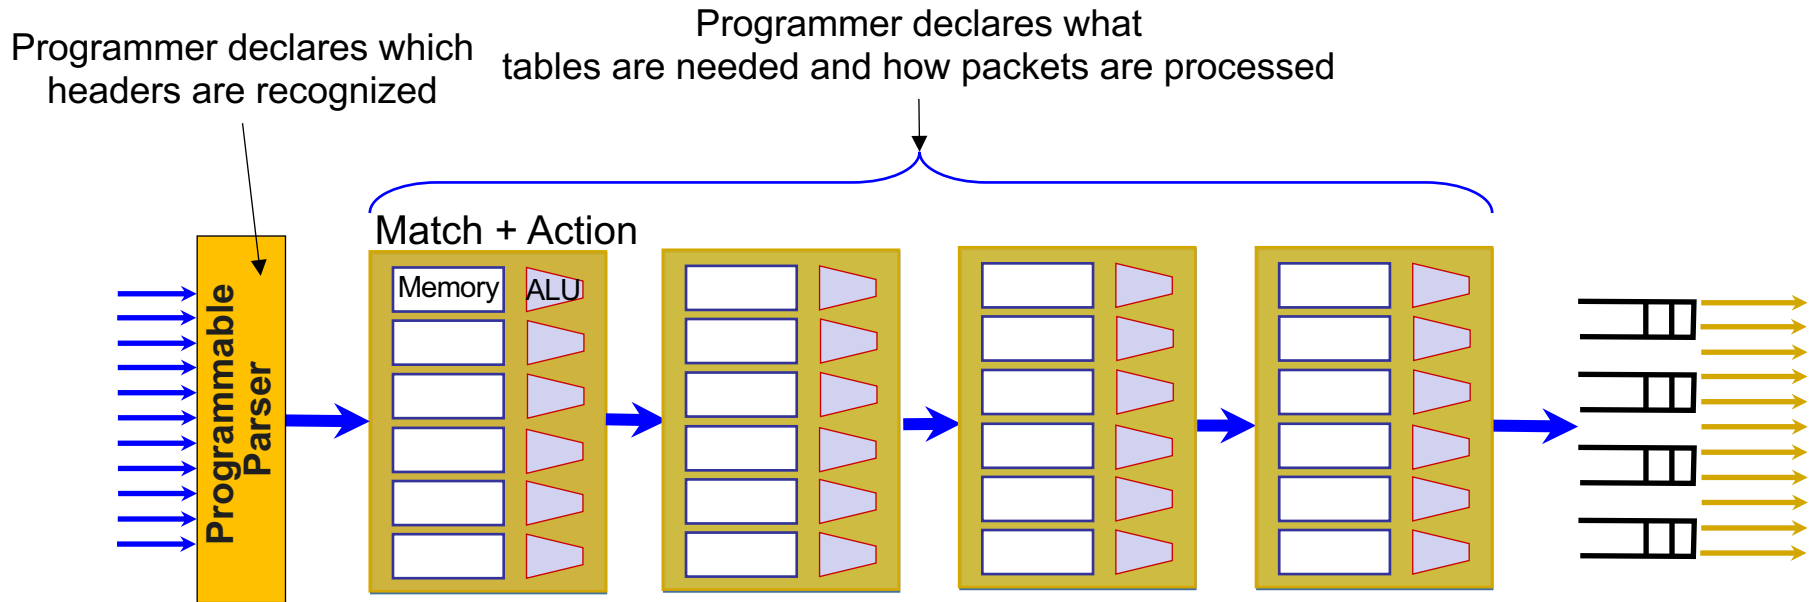
```
action export_queue_latency (sw_id) {  
    add_header(int_header);  
    modify_field(int_header.kind, TCP_OPTION_INT);  
    modify_field(int_header.len, TCP_OPTION_INT_LEN);  
    modify_field(int_header.sw_id, sw_id);  
    modify_field(int_header.q_latency,  
                intrinsic_metadata.deq_timedelta);  
    add_to_field(tcp.dataOffset, 2);  
    add_to_field(ipv4.totalLen, 8);  
    subtract_from_field(ingress_metadata.tcpLength, 12);  
}
```

# What's left?

---

- Compile it down to be something runnable on a programmable switch
  - Similar to other high-level languages we use to run code on hardware like CPU, GPU, FPGA etc.
  - P4 for programmable switches
- But which switch?

# PISA: Protocol Independent Switch Architecture



**All stages are identical – makes PISA a good “compiler target”**

# How's programmability used today?

---

- Remove features to reduce complexity
- Add proprietary features
- Silicon independence or avoid vendor lock-in
- Telemetry and measurements

# Example: In-band network telemetry (INT)

---

- “Which path did my packet take?”
- “Which rules did my packet follow?”
- “How long did it queue at each switch?”
- “Who did it share the queues with?”

# Why now?

---

- One of the earlier incarnation of programmable networks was in mid 90s
  - Active networks
- What's changed after two+ decades?
  - **Hardware:** We can now make programmable switches as fast as fixed ones
  - **Software:** We have found a (so far) reasonable balance between programmability, performance, and security

# Summary

---

- Abstractions beget modularity
  - Modularity is (almost always) good
- Programmability is powerful
  - Finding the right balance is hard
- Next lecture: Layer 2



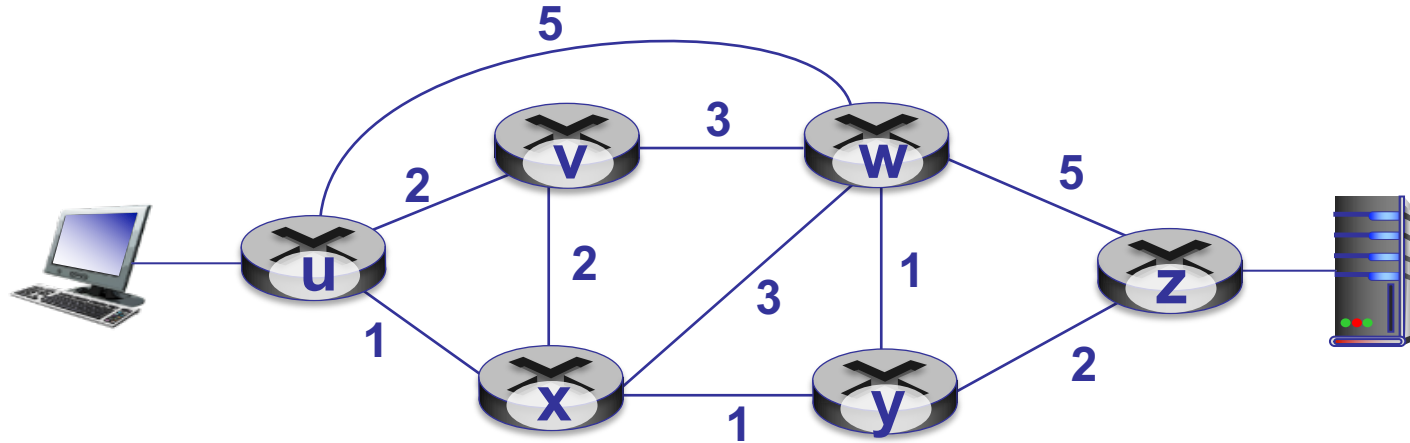


# Traffic engineering

---

- Want to avoid persistent overloads on links
- Choose routes to spread traffic load across links

# Traffic engineering: Difficult

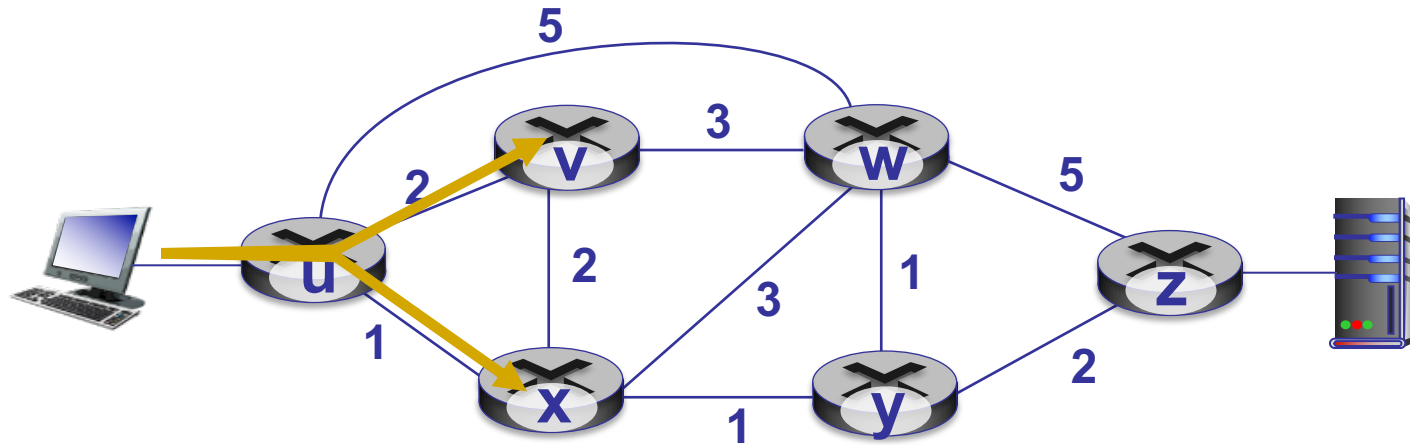


**Q:** What if network operator wants u-to-z traffic to flow along uvwz, x-to-z traffic to flow xwyz?

**A:** Need to define link weights so traffic routing algorithm computes routes accordingly (or need a new routing algorithm)!

Link weights are only control “knobs”

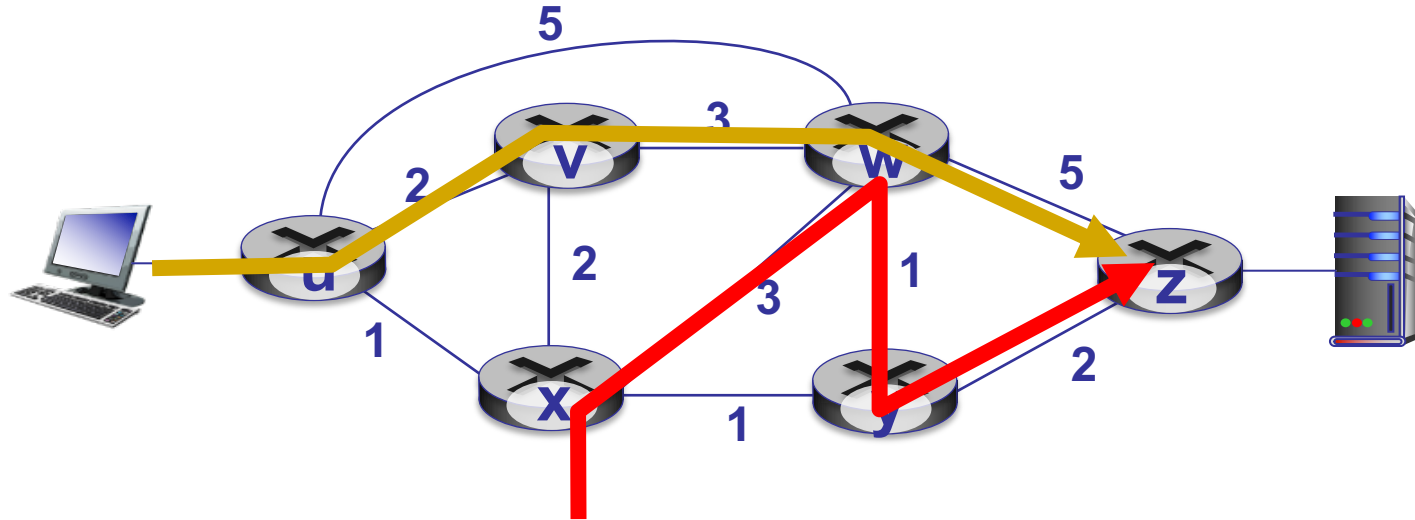
# Traffic engineering: Difficult



**Q:** What if network operator wants to split u-to-z traffic along uvwz and uxyz (load balancing)?

**A:** Can't do it (or need a new routing algorithm)

# Traffic engineering: Difficult



**Q:** What if w wants to route the two flows differently?

**A:** Can't do it (with LS or DV)

# OpenFlow: Controller-to-switch messages

---

- Key controller-to-switch messages
  - **Features**: controller queries switch features, switch replies
  - **Configure**: controller queries/sets switch configuration parameters
  - **Modify-state**: add, delete, modify flow entries in the OpenFlow tables
  - **Packet-out**: controller can send this packet out of specific switch port

# OpenFlow: Switch-to-controller messages

---

- Key switch-to-controller messages
  - **Packet-in**: transfer packet (and its control) to controller. See packet-out message from controller
  - **Flow-removed**: flow table entry deleted at switch
  - **Port status**: inform controller of a change on a port
- Network operators do not “program” switches by creating/sending OpenFlow messages directly.
  - Instead, they use higher-level abstraction at controller