

# **EECS 489**

# **Computer Networks**

**Fall 2019**

Mosharaf Chowdhury

*Material with thanks to Aditya Akella, Sugih Jamin, Philip Levis, Sylvia Ratnasamy, Peter Steenkiste, and many other colleagues.*

# Agenda

---

- CDN: Content Distribution Network
- DNS: Domain Name System

# Recap:

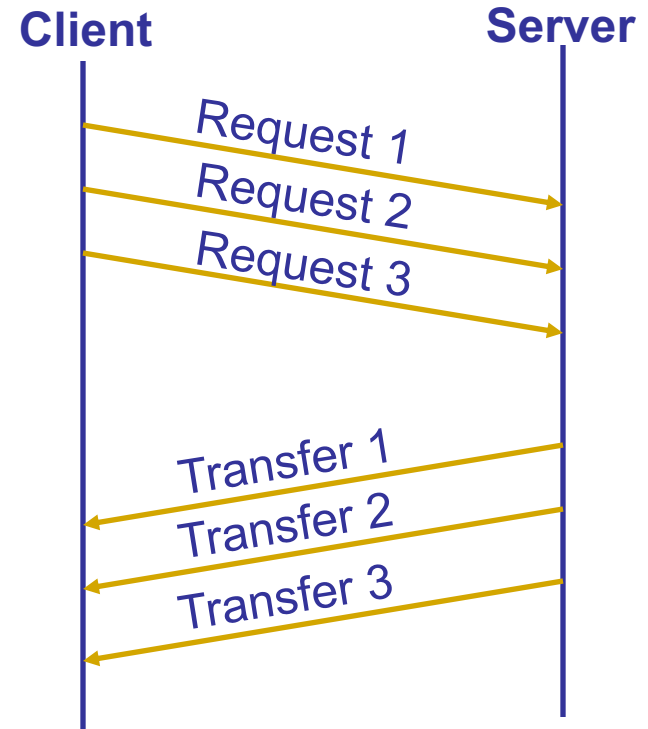
# Improving HTTP performance

---

- Optimizing connections using three “P”s
  - Persistent connections
  - Parallel/concurrent connections
  - Pipelined transfers over the same connection
- Caching
  - Forward proxy: close to clients
  - Reverse proxy: close to servers
- Replication

# Why pipeline and why not?

- Data are sent in a FIFO manner
  - Can lead to **head-of-line (HOL) blocking** if many small responses follow a large one
  - Not supported by default by major browsers circa 2015
- **Solution**
  - Priority and preemption
  - HTTP/2



# Replication

---

- Replicate popular Websites across many machines
  - Spreads load across servers
  - Places content closer to clients
  - Helps when content isn't cacheable

# Content Distribution Networks (CDN)

---

- Caching and replication as a service
- Large-scale distributed storage infrastructure (usually) administered by one entity
  - e.g., Akamai has servers in 20,000+ locations
- Combination of caching and replication
  - **Pull**: Direct result of clients' requests (caching)
  - **Push**: Expectation of high access rate (replication)
- Can do some processing to handle dynamic webpage content

# Cost-effective content delivery

---

- General theme: multiple sites hosted on shared physical infrastructure
  - Efficiency of statistical multiplexing
  - Economies of scale (volume pricing, etc.)
  - Amortization of human operator costs
- Examples:
  - CDNs
  - Web hosting companies
  - Cloud infrastructure

# CDN example – Akamai

---

- Akamai creates new domain names for each client
  - e.g., [a128.g.akamai.net](http://a128.g.akamai.net) for [cnn.com](http://cnn.com)
- The client content provider modifies content so that embedded URLs reference new domains
  - “Akamaize” content
  - e.g., <http://www.cnn.com/image-of-the-day.gif> becomes <http://a128.g.akamai.net/image-of-the-day.gif>
- Requests now sent to CDN’s infrastructure



# Why direct clients to particular replicas?

---

- Balancing load across server replicas
- Pairing clients with nearby servers to decrease latency and overall bandwidth usage

---

# **DNS: DOMAIN NAME SYSTEM**

# Internet names & addresses

---

- Machine addresses: e.g., 141.212.113.143
  - Router-usable labels for machines
  - Conforms to network structure (the “where”)
- Machine names: e.g., cse.umich.edu
  - Human-usable labels for machines
  - Conforms to organizational structure (the “who”)
- The Domain Name System (DNS) is how we map from one to the other
  - A **directory** service

# Why?

---

- Convenience
  - Easier to remember www.google.com than 216.58.216.100
- Provides a **level of indirection!**
  - Decoupled names from addresses
  - Many uses beyond just naming a specific host

# DNS: History

---

- Initially all host-address mappings were in a `hosts.txt` file (in `/etc/hosts`):
  - Maintained by the Stanford Research Institute (SRI)
  - Changes were submitted by email and updates downloaded periodically from SRI
- As the Internet grew SRI couldn't handle load
  - Names were not unique anymore
  - Hosts had inaccurate copies of `hosts.txt`

# Goals

---

- Uniqueness: no naming conflicts
- Scalable
  - Many names and frequent updates (secondary)
- Distributed, autonomous administration
  - Ability to update my own (machines') names
  - Don't have to track everybody's updates
- Highly available
- Lookups are fast
- Perfect consistency is a non-goal

# How?

---

- Partition the namespace
- Distribute administration of each partition
  - Autonomy to update my own (machines') names
  - Don't have to track everybody's updates
- Distribute name resolution for each partition
- How should we partition things?

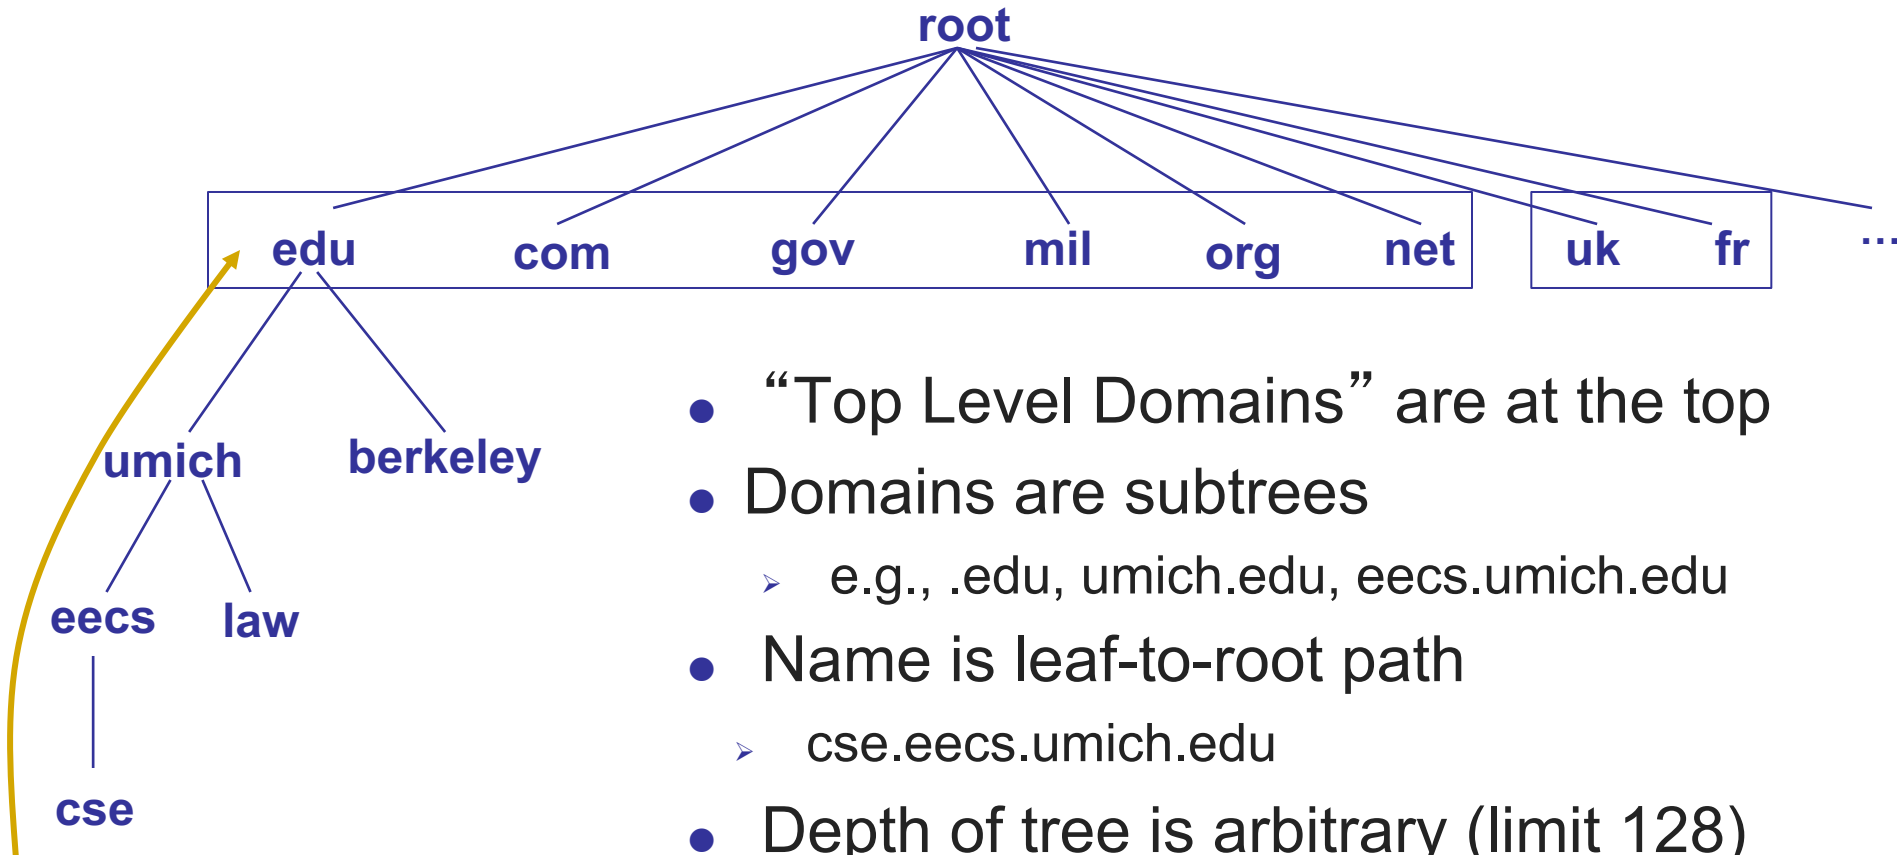
# Key idea: Hierarchy

---

- Three intertwined hierarchies
  - Hierarchical namespace
    - » As opposed to original flat namespace
  - Hierarchically administered
    - » As opposed to centralized
  - (Distributed) hierarchy of servers
    - » As opposed to centralized storage

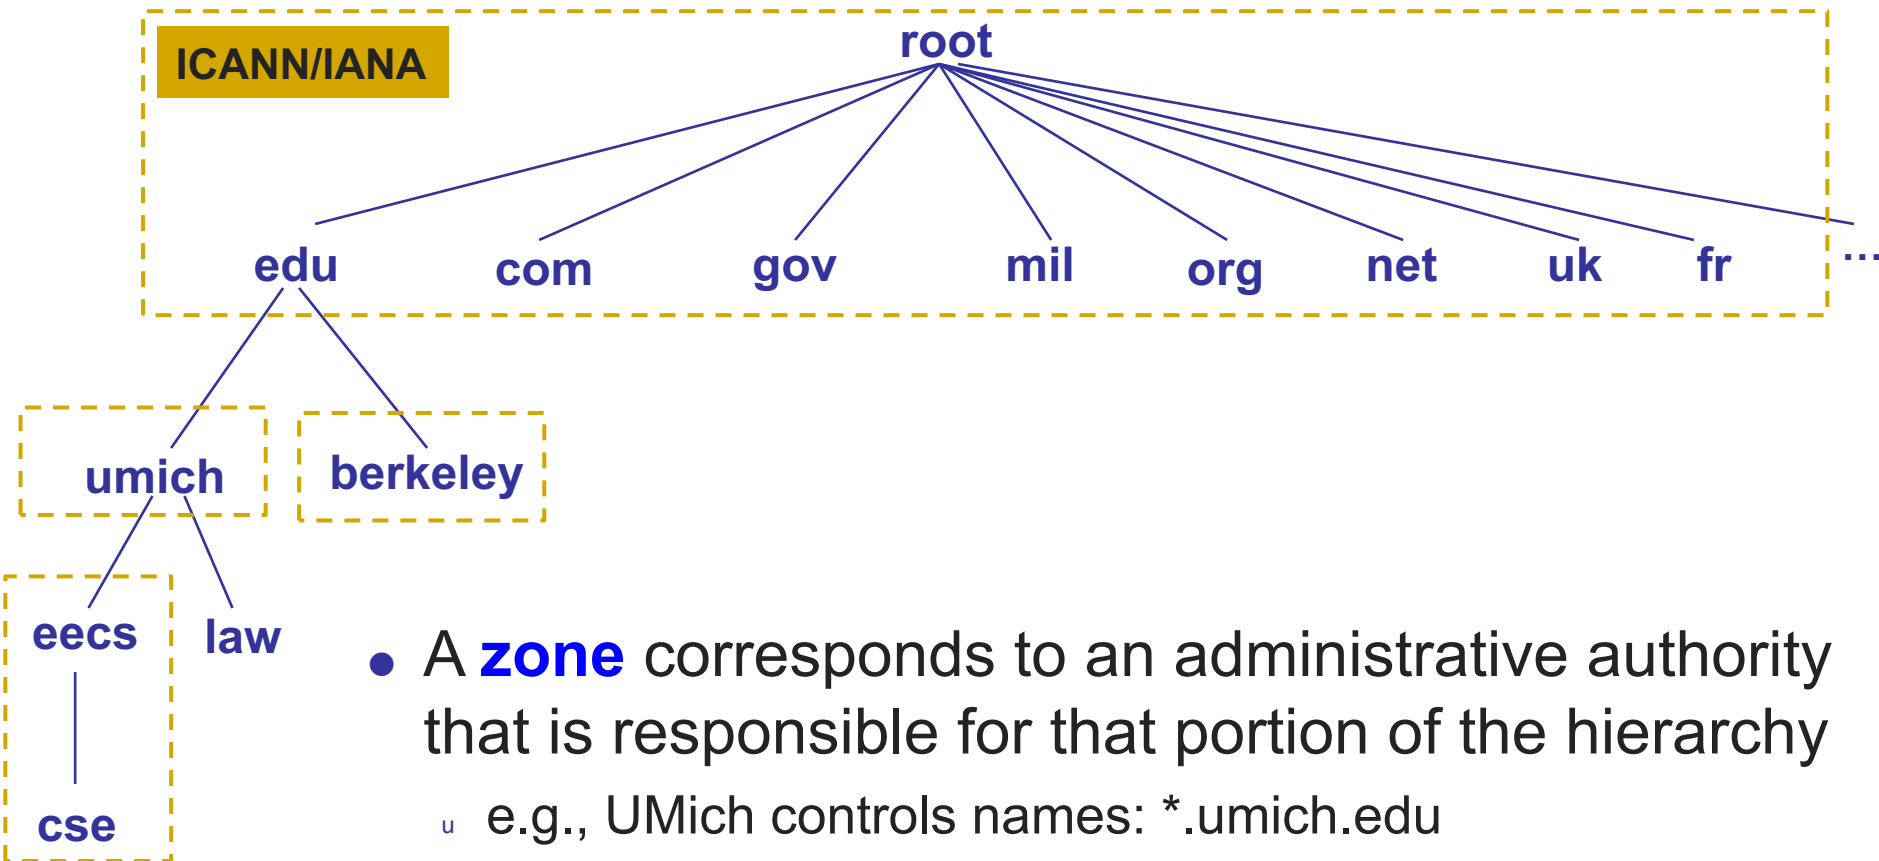


# Hierarchical namespace



- “Top Level Domains” are at the top
- Domains are subtrees
  - e.g., .edu, umich.edu, eeecs.umich.edu
- Name is leaf-to-root path
  - cse.eeecs.umich.edu
- Depth of tree is arbitrary (limit 128)
- Name collisions trivially avoided
  - Each domain is responsible

# Hierarchical administration



- A **zone** corresponds to an administrative authority that is responsible for that portion of the hierarchy
  - u e.g., UMich controls names: \*.umich.edu
  - u e.g., EECS controls names: \*.eeecs.umich.edu

# Server hierarchy

---

- Top of hierarchy: **Root servers**
  - Location hardwired into other servers
- Next Level: **Top-level domain (TLD) servers**
  - .com, .edu, etc.
  - Managed professionally
- Bottom Level: **Authoritative DNS servers**
  - Actually store the name-to-address mapping
  - Maintained by the corresponding administrative authority

# Server hierarchy

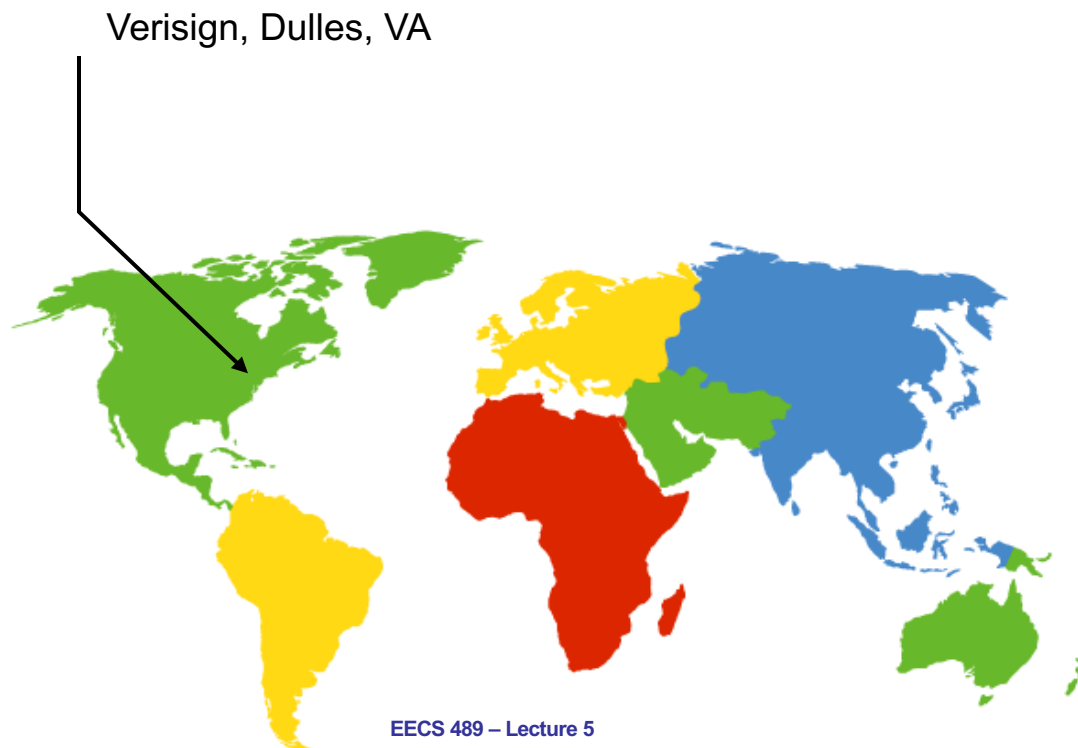
---

- Each server stores a (small!) subset of the total DNS database
- An authoritative DNS server stores “resource records” for all DNS names in the domain that it has authority for
- Each server needs to know other servers responsible for other portions of the hierarchy
  - Every server knows the root
  - Root server knows about all top-level domains

# DNS root

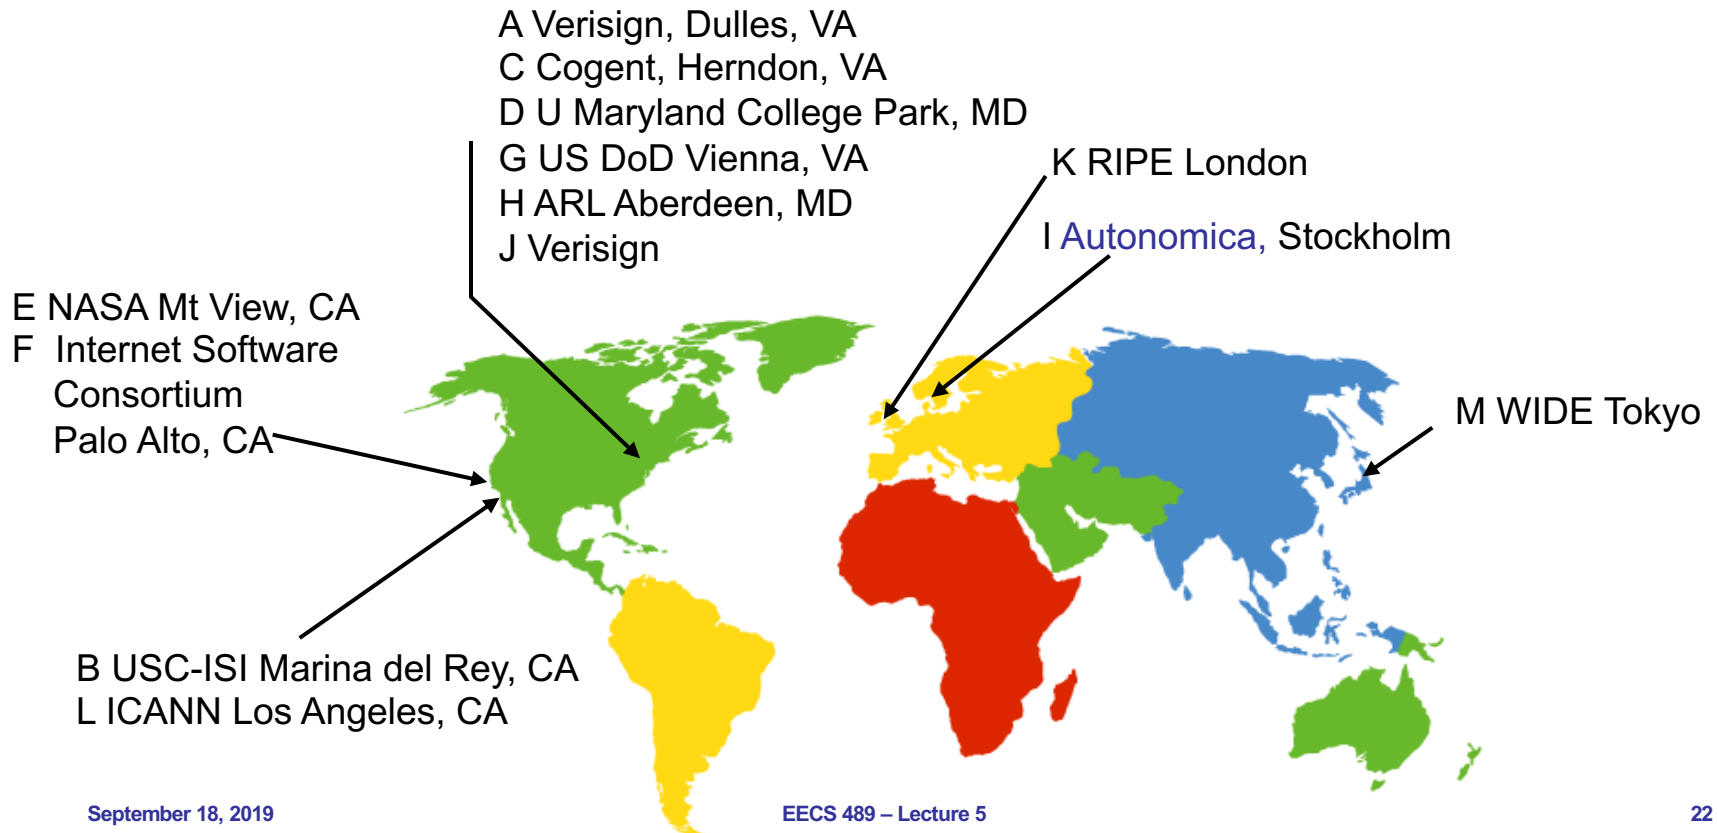
---

- Located in Virginia, USA
- How do we make the root scale?



# DNS root servers

- 13 root servers (labeled A-M; see <http://www.root-servers.org/>)



---

**5-MINUTE BREAK!**

# Announcements

---

- **Assignment 1 due on September 26**
  - Autograder out soon
- TCP is provides stream-level guarantees
  - » Not message-level



# DNS records

---

- DNS servers store **resource records (RRs)**
  - RR is (name, value, type, TTL)
- Type = A: (→ **A**ddress)
  - name = hostname
  - value = IP address
- Type = NS: (→ **N**ame **S**erver)
  - name = domain
  - value = name of DNS server for domain

# DNS records (cont'd)

---

- Type = CNAME: (→ Canonical Name)
  - name = alias name for some “canonical” (real) name
    - » e.g., cse.umich.edu is really cse.eecs.umich.edu
  - value = canonical name
- Type = MX: (→ Mail eXchanger)
  - name = domain in email address
  - value = name(s) of mail server(s)

# Inserting Resource Records into DNS

---

- Register foobar.com at registrar
  - Provide registrar with names and IP addresses of your authoritative name server(s)
  - Registrar inserts RR pairs into the .com TLD server:
    - »(foobar.com, dns1.foobar.com, NS)
    - »(dns1.foobar.com, 212.44.9.129, A)
- Store resource records in your server dns1.foobar.com
  - e.g., type A record for www.foobar.com
  - e.g., type MX record for foobar.com


# Using DNS (Client/App View)

---

- Two components
  - Local DNS servers
  - Resolver software on hosts
- Local DNS server (“default name server”)
  - Clients configured with default server’s address OR learn it via a host configuration protocol (e.g., DHCP)
- Client application
  - Obtain DNS name (e.g., from URL)
  - Do `gethostbyname()` to trigger DNS request to its local DNS server


# Name resolution

---




# Name resolution

---




# Name resolution

---



# Name resolution


---






# Name resolution: Recursive

---



# Name resolution: Iterative



# Two ways to resolve a name

---

- Recursive name resolution
  - Ask server to do it for you
- Iterative name resolution
  - Ask server who to ask next
- The iterative example we saw is a mix of both!

# DNS protocol

---

- Query and Reply messages; both with the same message format
  - Header: identifier, flags, etc.
  - Plus resource records
- Client–server interaction on UDP Port 53
  - Spec supports TCP too, but not always implemented

# Goals: Are we there yet?

---

- Uniqueness: No naming conflicts
- Scalable
- Distributed, autonomous administration
- Highly available?

# Reliability

---

- **Replicated** DNS servers (primary/secondary)
  - Name service available if at least one replica is up
  - Queries can be load-balanced between replicas
- Usually, UDP used for queries
  - Reliability, if needed, must be implemented on UDP
- Try alternate servers on timeout
  - **Exponential backoff** when retrying same server
- Same identifier for all queries
  - Don't care which server responds

# Goals: Are we there yet?

---

- Uniqueness: No naming conflicts
- Scalable
- Distributed, autonomous administration
- Highly available
- Fast lookups?

# DNS caching

---

- Performing all these queries takes time
  - Up to 1-second latency before starting download
- Caching can greatly reduce overhead
  - The top-level servers very rarely change
  - Popular sites (e.g., `www.cnn.com`) visited often
  - Local DNS server often has the information cached
- How DNS caching works
  - DNS servers cache responses to queries
  - Responses include a “time to live” (TTL) field
  - Server deletes cached entry after TTL expires



# Negative caching

---

- Remember things that don't work
  - Misspellings like `www.cnn.comm` and `www.cnnn.com`
  - These can take a long time to fail the first time
  - Good to remember that they don't work so the failure takes less time the next time around
- Negative caching is optional
  - Not widely implemented

# Important properties of DNS

---

- Administrative delegation and hierarchy enables:
  - Easy unique naming
  - “Fate sharing” for network failures
  - Reasonable trust model
  - Caching increases scalability and performance

# DNS provides indirection

---

- Addresses can change underneath
  - Move `www.cnn.com` to `4.125.91.21`
- Name could map to multiple IP addresses
  - Load-balancing (CDN)
  - Reducing latency by picking nearby servers (CDN)
- Multiple names for the same address
  - E.g., many services (mail, www) on same machine
  - E.g., aliases like `www.cnn.com` and `cnn.com`
- This flexibility applies only within domain!

# Summary

---

- CDNs improve web performance
  - Via replication and caching
  - Good server selection
- DNS allows us to go to webpages without having to memorize IP addresses
  - Allows a level of indirection that enables many functionalities including CDN server selection